

Final_Project

Libraries

```
pacman::p_load(data.table,tidyverse,magrittr,YARF,skimr,plyr,tidyr,YARF,mltools,caret)
```

Loading Data


```
df= read.csv("housing_data_2016_2017.csv")
head(df,1)
```

```
##                               HITId                               HITTypeId
## 1 30ID399FXG7F26JWONXFOY86J90FD4 36BILMLQB75QQNBTKGYCZWDN8TVAU
##                               Title
## 1 Find Information about Housing To Help a Student Project -- Very easy
##                               Description Keywords Reward
## 1 Go to a link and copy information into the HIT          NA $0.05
##                               CreationTime MaxAssignments
## 1 Wed Feb 15 22:13:37 PST 2017              1
##                               RequesterAnnotation AssignmentDurationInSeconds
## 1 BatchId:2689947;OriginalHitTemplateId:920937336;              900
##   AutoApprovalDelayInSeconds              Expiration NumberOfSimilarHITs
## 1              60 Wed Feb 22 22:13:37 PST 2017              NA
##   LifetimeInSeconds              AssignmentId              WorkerId
## 1              NA 32KTQ2V7RDFCSAWQOW1SXC5AZIC9MB A231MNJJDDF3LS
##   AssignmentStatus              AcceptTime              SubmitTime
## 1   Approved Thu Feb 16 05:32:36 PST 2017 Thu Feb 16 05:35:37 PST 2017
##   AutoApprovalTime              ApprovalTime RejectionTime
## 1 Thu Feb 16 05:36:37 PST 2017 2017-02-16 13:37:11 UTC          NA
##   RequesterFeedback WorkTimeInSeconds LifetimeApprovalRate
## 1              NA              181              100% (187/187)
##   Last30DaysApprovalRate Last7DaysApprovalRate
## 1              100% (187/187)              100% (187/187)
##
##                               URL
## 1 http://www.mlsli.com/homes-for-sale/address-not-available-from-broker-Flushing-NY-11355-149238320
##   approx_year_built cats_allowed common_charges community_district_num
## 1              1955              no              $767              25
##   coop_condo date_of_sale dining_room_type dogs_allowed fuel_type
## 1   co-op   2/16/2016              combo              no              gas
##   full_address_or_zip_code garage_exists kitchen_type maintenance_cost
## 1   Flushing NY, 11355              <NA>              eat in              <NA>
##   model_type num_bedrooms num_floors_in_building num_full_bathrooms
## 1 Mitchell Garden 3              2              6              1
##   num_half_bathrooms num_total_rooms parking_charges pct_tax_deductibl
## 1              NA              5              <NA>              NA
##   sale_price sq_footage total_taxes walk_score listing_price_to_nearest_1000
## 1   $228,000              NA              <NA>              82              <NA>
##   url
```

```
## 1 <NA>
```

```
hist(as.numeric(gsub('[$,]', '', as.character(df$sale_price))))
```

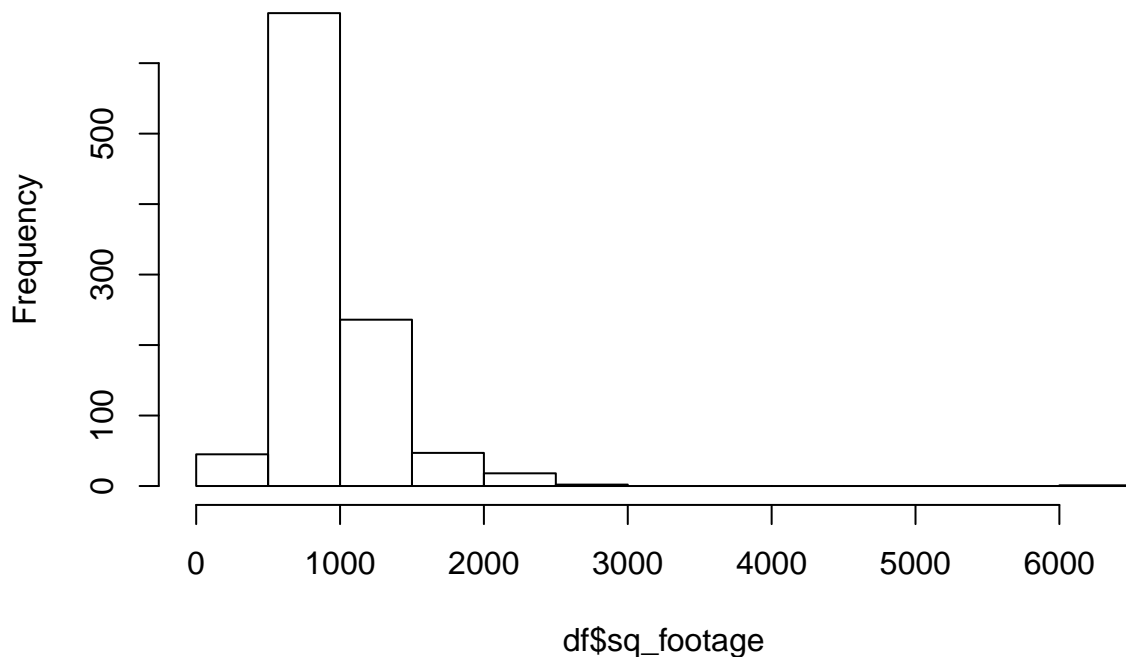
Histogram of `as.numeric(gsub("[$,]", "", as.character(df$sale_price))`



`as.numeric(gsub("[$,]", "", as.character(df$sale_price)))`

```
hist(df$sq_footage)
```

Histogram of `df$sq_footage`



##

Useful Summary of Data This gives a broad overview on the data,

```
skim(df)
```

Table 1: Data summary

Name	df
Number of rows	2230
Number of columns	55
Column type frequency:	
factor	36
logical	5
numeric	14
Group variables	None

Variable type: factor

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
HITId	758	0.66	FALSE	1472	301: 1, 301: 1, 301: 1, 302: 1
HITTypeId	758	0.66	FALSE	2	310: 944, 36B: 528
Title	758	0.66	FALSE	1	Fin: 1472
Description	758	0.66	FALSE	2	Got: 944, Go : 528
Reward	758	0.66	FALSE	1	\$0.: 1472
CreationTime	758	0.66	FALSE	62	Thu: 43, Thu: 40, Wed: 39, Thu: 37
RequesterAnnotation	758	0.66	FALSE	2	Bat: 944, Bat: 528
Expiration	758	0.66	FALSE	62	Thu: 43, Thu: 40, Wed: 39, Thu: 37
AssignmentId	758	0.66	FALSE	1472	301: 1, 301: 1, 304: 1, 304: 1
WorkerId	758	0.66	FALSE	73	A23: 187, A1S: 129, A3C: 124, AHX
AssignmentStatus	758	0.66	FALSE	1	App: 1472
AcceptTime	758	0.66	FALSE	1457	Thu: 2, Thu: 2, Thu: 2, Thu: 2
SubmitTime	758	0.66	FALSE	1460	Thu: 2, Thu: 2, Thu: 2, Thu: 2
AutoApprovalTime	758	0.66	FALSE	1460	Thu: 2, Thu: 2, Thu: 2, Thu: 2
ApprovalTime	758	0.66	FALSE	929	201: 6, 201: 6, 201: 5, 201: 5
LifetimeApprovalRate	758	0.66	FALSE	32	100: 187, 100: 126, 100: 124, 100: 10
Last30DaysApprovalRate	758	0.66	FALSE	32	100: 187, 100: 126, 100: 124, 100: 10
Last7DaysApprovalRate	758	0.66	FALSE	32	100: 187, 100: 126, 100: 124, 100: 10
URL	758	0.66	FALSE	1450	htt: 2, htt: 2, htt: 2, htt: 2
cats_allowed	0	1.00	FALSE	3	no: 1402, yes: 826, y: 2
common_charges	1684	0.24	FALSE	258	\$25: 11, \$17: 10, \$27: 9, \$29: 8
coop_condo	0	1.00	FALSE	2	co-: 1661, con: 569
date_of_sale	1702	0.24	FALSE	222	6/3: 7, 10/: 6, 12/: 6, 2/2: 6
dining_room_type	448	0.80	FALSE	5	com: 957, for: 620, oth: 201, din: 2
dogs_allowed	0	1.00	FALSE	3	no: 1684, yes: 544, yes: 2
fuel_type	112	0.95	FALSE	6	gas: 1348, oil: 664, ele: 62, oth: 40
full_address_or_zip_code	0	1.00	FALSE	1177	70-: 22, 269: 17, 270: 16, 73-: 14
garage_exists	1826	0.18	FALSE	6	yes: 361, Yes: 39, 1: 1, eys: 1
kitchen_type	16	0.99	FALSE	13	eat: 733, eff: 505, com: 349, eff: 338
maintenance_cost	623	0.72	FALSE	609	\$54: 10, \$67: 10, \$68: 10, \$70: 10
model_type	40	0.98	FALSE	875	1 B: 63, One: 59, 2 B: 50, Hi-: 41
parking_charges	1671	0.25	FALSE	89	\$15: 42, \$60: 41, \$75: 27, \$13: 23
sale_price	1702	0.24	FALSE	315	\$15: 11, \$17: 10, \$13: 7, \$22: 7

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
total_taxes	1646	0.26	FALSE	293	\$13: 13, \$25: 12, \$4,: 11, \$2,: 10
listing_price_to_nearest_1000	534	0.76	FALSE	292	\$34: 28, \$39: 26, \$28: 25, \$23: 23
url	758	0.66	FALSE	1450	htt: 2, htt: 2, htt: 2, htt: 2

Variable type: logical

skim_variable	n_missing	complete_rate	mean	count
Keywords	2230	0	NaN	:
NumberOfSimilarHITs	2230	0	NaN	:
LifetimeInSeconds	2230	0	NaN	:
RejectionTime	2230	0	NaN	:
RequesterFeedback	2230	0	NaN	:

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
MaxAssignments	758	0.66	1.00	0.00	1	1	1	1	1	
AssignmentDurationInSeconds	758	0.66	900.00	0.00	900	900	900	900	900	
AutoApprovalDelayInSeconds	758	0.66	60.00	0.00	60	60	60	60	60	
WorkTimeInSeconds	758	0.66	162.39	111.69	22	89	127	197	815	
approx_year_built	40	0.98	1962.71	21.08	1893	1950	1958	1970	2017	
community_district_num	19	0.99	26.33	2.95	3	25	26	28	32	
num_bedrooms	115	0.95	1.65	0.74	0	1	2	2	6	
num_floors_in_building	650	0.71	7.79	7.52	1	3	6	7	34	
num_full_bathrooms	0	1.00	1.23	0.44	1	1	1	1	3	
num_half_bathrooms	2058	0.08	0.95	0.30	0	1	1	1	2	
num_total_rooms	2	1.00	4.14	1.35	0	3	4	5	14	
pct_tax_deductibl	1754	0.21	45.40	6.95	20	40	50	50	75	
sq_footage	1210	0.46	955.36	380.86	100	743	881	1100	6215	
walk_score	0	1.00	83.92	14.75	7	77	89	95	99	

There is a lot of data that is completely missing and some that is heavily missing. I decided to remove them. Some examples below.

Keywords,NumberOfSimilarHITs, LifetimeInSeconds, RejectionTime,RequesterFeedback all completely missing. ommon_charges(missing 1684),garage_exists(missing 1826)

```
cat("Data has",nrow(df),"number of rows\n")
```

```
## Data has 2230 number of rows
```

```
cat("Data has",ncol(df), "number of columns")
```

```
## Data has 55 number of columns
```

```
sort(colMeans(is.na(df)), decreasing = TRUE)
```

```
##           Keywords           NumberOfSimilarHITs
##           1.000000000           1.000000000
##           LifetimeInSeconds           RejectionTime
```

##	1.000000000	1.000000000
##	RequesterFeedback	num_half_bathrooms
##	1.000000000	0.922869955
##	garage_exists	pct_tax_deductibl
##	0.818834081	0.786547085
##	date_of_sale	sale_price
##	0.763228700	0.763228700
##	common_charges	parking_charges
##	0.755156951	0.749327354
##	total_taxes	url
##	0.738116592	0.660089686
##	sq_footage	HITId
##	0.542600897	0.339910314
##	HITTypeId	Title
##	0.339910314	0.339910314
##	Description	Reward
##	0.339910314	0.339910314
##	CreationTime	MaxAssignments
##	0.339910314	0.339910314
##	RequesterAnnotation	AssignmentDurationInSeconds
##	0.339910314	0.339910314
##	AutoApprovalDelayInSeconds	Expiration
##	0.339910314	0.339910314
##	AssignmentId	WorkerId
##	0.339910314	0.339910314
##	AssignmentStatus	AcceptTime
##	0.339910314	0.339910314
##	SubmitTime	AutoApprovalTime
##	0.339910314	0.339910314
##	ApprovalTime	WorkTimeInSeconds
##	0.339910314	0.339910314
##	LifetimeApprovalRate	Last30DaysApprovalRate
##	0.339910314	0.339910314
##	Last7DaysApprovalRate	URL
##	0.339910314	0.339910314
##	num_floors_in_building	maintenance_cost
##	0.291479821	0.279372197
##	listing_price_to_nearest_1000	dining_room_type
##	0.239461883	0.200896861
##	num_bedrooms	fuel_type
##	0.051569507	0.050224215
##	approx_year_built	model_type
##	0.017937220	0.017937220
##	community_district_num	kitchen_type
##	0.008520179	0.007174888
##	num_total_rooms	cats_allowed
##	0.000896861	0.000000000
##	coop_condo	dogs_allowed
##	0.000000000	0.000000000
##	full_address_or_zip_code	num_full_bathrooms
##	0.000000000	0.000000000
##	walk_score	
##	0.000000000	

Data Cleaning Remove all missing y

```
df_drops = df %>% drop_na(sale_price)
# skim(df_drops) %>%
#   summary()
skim(df_drops)
```

Table 5: Data summary

Name	df_drops
Number of rows	528
Number of columns	55
Column type frequency:	
factor	36
logical	5
numeric	14
Group variables	None

Variable type: factor

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
HITId	0	1.00	FALSE	528	301: 1, 302: 1, 302: 1, 307: 1
HITTypeId	0	1.00	FALSE	1	36B: 528, 310: 0
Title	0	1.00	FALSE	1	Fin: 528
Description	0	1.00	FALSE	1	Go : 528, Got: 0
Reward	0	1.00	FALSE	1	\$0.: 528
CreationTime	0	1.00	FALSE	21	Wed: 39, Wed: 36, Wed: 33, Wed: 3
RequesterAnnotation	0	1.00	FALSE	1	Bat: 528, Bat: 0
Expiration	0	1.00	FALSE	21	Wed: 39, Wed: 36, Wed: 33, Wed: 3
AssignmentId	0	1.00	FALSE	528	301: 1, 301: 1, 308: 1, 308: 1
WorkerId	0	1.00	FALSE	21	A23: 187, AHX: 102, A1K: 80, A3S:
AssignmentStatus	0	1.00	FALSE	1	App: 528
AcceptTime	0	1.00	FALSE	523	Thu: 2, Thu: 2, Thu: 2, Thu: 2
SubmitTime	0	1.00	FALSE	524	Thu: 2, Thu: 2, Thu: 2, Thu: 2
AutoApprovalTime	0	1.00	FALSE	524	Thu: 2, Thu: 2, Thu: 2, Thu: 2
ApprovalTime	0	1.00	FALSE	337	201: 5, 201: 5, 201: 4, 201: 4
LifetimeApprovalRate	0	1.00	FALSE	15	100: 187, 100: 102, 100: 80, 100: 58
Last30DaysApprovalRate	0	1.00	FALSE	15	100: 187, 100: 102, 100: 80, 100: 58
Last7DaysApprovalRate	0	1.00	FALSE	15	100: 187, 100: 102, 100: 80, 100: 58
URL	0	1.00	FALSE	524	htt: 2, htt: 2, htt: 2, htt: 2
cats_allowed	0	1.00	FALSE	2	no: 285, yes: 243, y: 0
common_charges	396	0.25	FALSE	112	\$27: 3, \$31: 3, \$21: 2, \$24: 2
coop_condo	0	1.00	FALSE	2	co-: 399, con: 129
date_of_sale	0	1.00	FALSE	222	6/3: 7, 10/: 6, 12/: 6, 2/2: 6
dining_room_type	120	0.77	FALSE	4	com: 241, for: 116, oth: 49, din: 2
dogs_allowed	0	1.00	FALSE	2	no: 381, yes: 147, yes: 0
fuel_type	24	0.95	FALSE	6	gas: 301, oil: 180, ele: 11, oth: 8
full_address_or_zip_code	0	1.00	FALSE	468	70-: 8, 54-: 4, 104: 3, 117: 3
garage_exists	434	0.18	FALSE	6	yes: 51, Yes: 39, 1: 1, eys: 1
kitchen_type	6	0.99	FALSE	7	eff: 231, eat: 190, Com: 50, com: 31

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
maintenance_cost	142	0.73	FALSE	284	\$52: 4, \$60: 4, \$66: 4, \$67: 4
model_type	15	0.97	FALSE	356	1 B: 23, One: 19, 2 B: 11, Gar: 11
parking_charges	393	0.26	FALSE	50	\$10: 12, \$20: 10, \$95: 8, \$12: 7
sale_price	0	1.00	FALSE	315	\$15: 11, \$17: 10, \$13: 7, \$22: 7
total_taxes	397	0.25	FALSE	120	\$2,: 3, \$4,: 3, \$1,: 2, \$1,: 2
listing_price_to_nearest_1000	528	0.00	FALSE	0	\$1,: 0, \$10: 0, \$10: 0, \$10: 0
url	0	1.00	FALSE	524	htt: 2, htt: 2, htt: 2, htt: 2

Variable type: logical

skim_variable	n_missing	complete_rate	mean	count
Keywords	528	0	NaN	:
NumberOfSimilarHITs	528	0	NaN	:
LifetimeInSeconds	528	0	NaN	:
RejectionTime	528	0	NaN	:
RequesterFeedback	528	0	NaN	:

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
MaxAssignments	0	1.00	1.00	0.00	1	1	1	1	1	
AssignmentDurationInSeconds	0	1.00	900.00	0.00	900	900	900	900	900	
AutoApprovalDelayInSeconds	0	1.00	60.00	0.00	60	60	60	60	60	
WorkTimeInSeconds	0	1.00	150.38	96.99	52	97	123	163	815	
approx_year_built	6	0.99	1962.38	20.56	1915	1950	1957	1968	2016	
community_district_num	1	1.00	26.30	2.99	3	25	26	28	30	
num_bedrooms	0	1.00	1.54	0.75	0	1	1	2	3	
num_floors_in_building	108	0.80	7.08	6.83	1	2	6	7	34	
num_full_bathrooms	0	1.00	1.20	0.42	1	1	1	1	3	
num_half_bathrooms	498	0.06	1.03	0.18	1	1	1	1	2	
num_total_rooms	0	1.00	4.02	1.20	1	3	4	5	8	
pct_tax_deductibl	429	0.19	44.99	8.09	20	40	50	50	65	
sq_footage	315	0.40	965.28	490.42	375	750	874	1010	6215	
walk_score	0	1.00	83.10	13.09	15	76	85	94	99	

Meaningful Features Data Cleaning

Finding meaningful features. These are features I believe are meaningful. `df_mutated` has all the features that I will be using. I am not looking at what's missing yet or how the data looks like, just looking for features that would be best to predict sales price.

```
colnames(df)
```

```
## [1] "HITId"           "HITTypeId"
## [3] "Title"           "Description"
## [5] "Keywords"        "Reward"
## [7] "CreationTime"    "MaxAssignments"
## [9] "RequesterAnnotation" "AssignmentDurationInSeconds"
## [11] "AutoApprovalDelayInSeconds" "Expiration"
```

```
## [13] "NumberOfSimilarHITS"      "LifetimeInSeconds"
## [15] "AssignmentId"             "WorkerId"
## [17] "AssignmentStatus"         "AcceptTime"
## [19] "SubmitTime"               "AutoApprovalTime"
## [21] "ApprovalTime"             "RejectionTime"
## [23] "RequesterFeedback"        "WorkTimeInSeconds"
## [25] "LifetimeApprovalRate"     "Last30DaysApprovalRate"
## [27] "Last7DaysApprovalRate"    "URL"
## [29] "approx_year_built"        "cats_allowed"
## [31] "common_charges"           "community_district_num"
## [33] "coop_condo"               "date_of_sale"
## [35] "dining_room_type"         "dogs_allowed"
## [37] "fuel_type"                 "full_address_or_zip_code"
## [39] "garage_exists"            "kitchen_type"
## [41] "maintenance_cost"         "model_type"
## [43] "num_bedrooms"             "num_floors_in_building"
## [45] "num_full_bathrooms"       "num_half_bathrooms"
## [47] "num_total_rooms"          "parking_charges"
## [49] "pct_tax_deductibl"        "sale_price"
## [51] "sq_footage"               "total_taxes"
## [53] "walk_score"               "listing_price_to_nearest_1000"
## [55] "url"
```

```
df_mutated = copy(df_drops)
```

```
df_mutated %<>%
```

```
  select(cats_allowed,common_charges,coop_condo,dining_room_type,dogs_allowed,fuel_type,garage_exists,m
sort(colMeans(is.na(df_mutated)), decreasing = TRUE)
```

```
##          garage_exists      total_taxes      common_charges
##          0.821969697          0.751893939          0.750000000
##          sq_footage      maintenance_cost      dining_room_type
##          0.596590909          0.268939394          0.227272727
## num_floors_in_building      fuel_type      model_type
##          0.204545455          0.045454545          0.028409091
##      approx_year_built community_district_num      cats_allowed
##          0.011363636          0.001893939          0.000000000
##          coop_condo      dogs_allowed      num_bedrooms
##          0.000000000          0.000000000          0.000000000
##      num_full_bathrooms      num_total_rooms      sale_price
##          0.000000000          0.000000000          0.000000000
##          walk_score
##          0.000000000
```

Feature Data Cleaning

I am now looking more closely to the data. Looking at this there are too many types of model_types 875 different times from original data with NA sale price this seems difficult to deal with so I will remove this. I discarded data with more than 50% of missingness.

```
df_mutated_features = copy(df_mutated)
```

```
df_mutated_features %<>%
```

```
  select(-model_type,-total_taxes, -community_district_num)#,-common_charges,-sq_footage)
skim(df_mutated_features) %>%
  summary()
```


Table 9: Data summary

Name	df_mutated_features
Number of rows	528
Number of columns	16
Column type frequency:	
factor	9
numeric	7
Group variables	None

```
sort(colMeans(is.na(df_mutated_features)), decreasing = TRUE)
```

```
##      garage_exists      common_charges      sq_footage
##      0.82196970      0.75000000      0.59659091
##      maintenance_cost      dining_room_type num_floors_in_building
##      0.26893939      0.22727273      0.20454545
##      fuel_type      approx_year_built      cats_allowed
##      0.04545455      0.01136364      0.00000000
##      coop_condo      dogs_allowed      num_bedrooms
##      0.00000000      0.00000000      0.00000000
##      num_full_bathrooms      num_total_rooms      sale_price
##      0.00000000      0.00000000      0.00000000
##      walk_score
##      0.00000000
```

Oberservations Data Cleaning

I am okay with the number of features I have now. Now Ill be cleaning the observations.

```
df_clean = copy(df_mutated_features)
```

```
# Fixing y to be just yes and reducing factors to just yes and no.
```

```
df_clean %<>%
```

```
  mutate(cats_allowed = as.factor(ifelse(cats_allowed == 'y' | cats_allowed == 'yes', 'yes', 'no'))) %>%
```

```
#Fixing yes89 to just yes and reducing factors to just yes and no
```

```
  mutate(dogs_allowed = as.factor(ifelse(dogs_allowed == 'yes89' | dogs_allowed == 'yes', 'yes', 'no'))) %>%
```

```
#mutate(sale_price = as.numeric(gsub('[$', '', as.character(df_clean$sale_price))))
```

```
mutate(sale_price = as.numeric(gsub('[$', '', as.character(df_clean$sale_price))) )%>%
```

```
mutate(common_charges = as.numeric(gsub('[$', '', as.character(df_clean$common_charges)))) %>%
```

```
mutate(maintenance_cost = as.numeric(gsub('[$', '', as.character(df_clean$maintenance_cost)))) %>%
```

```
mutate(garage_exists = ifelse(is.na(garage_exists), 0, 1))
```

```
#mutate(fuel_type = if(is.na(fuel_type)){fuel_type = 'other'})
```

```

#Very annoying this best way I found to combine two factor lvels

library(forcats)
df_clean$fuel_type = fct_collapse(df_clean$fuel_type, other = c("other", "Other"))
#df_clean$dining_room_type = fct_collapse(df_clean$dining_room_type, other= c('other', 'none', 'dining ar

# df_clean_sub = copy(df_clean)
# df_clean_sub = df_clean_sub[df_clean_sub$sale_price < 700000,]
#
# df_clean = df_clean_sub
#df_clean = df_clean[df_clean$sq_footage < 2500,]
#which(df_clean$sq_footage > 2500)

#df_clean = df_clean[-136,]

options(scipen=999)
max(df_clean$sq_footage, na.rm = TRUE)

## [1] 6215
min(df_clean$sq_footage, na.rm = TRUE)

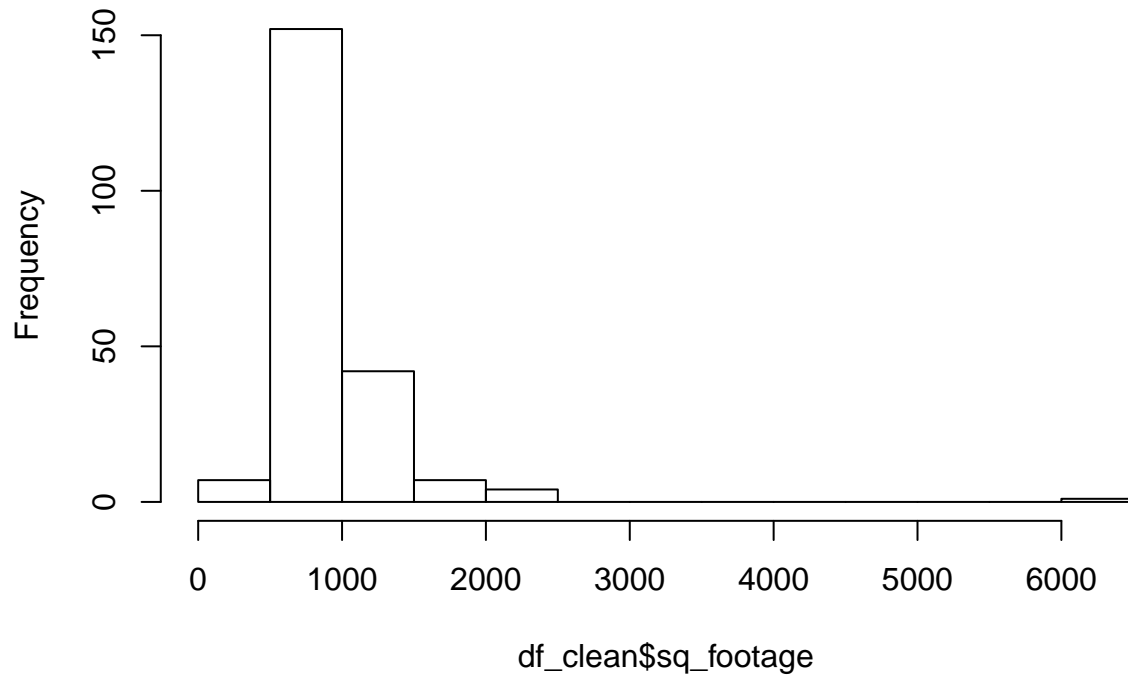
## [1] 375
max(df_clean$sale_price, na.rm = TRUE)

## [1] 999999
min(df_clean$sale_price, na.rm = TRUE)

## [1] 55000
hist(df_clean$sq_footage)

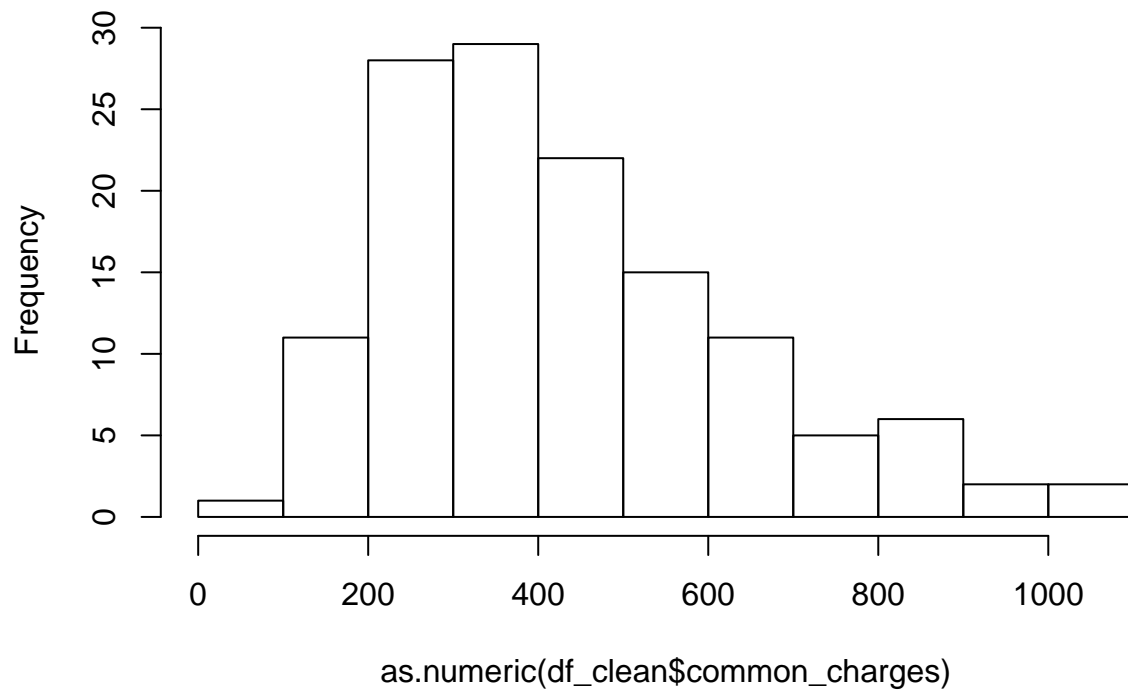
```

Histogram of df_clean\$sq_footage




```
hist(as.numeric(df_clean$common_charges))
```

Histogram of as.numeric(df_clean\$common_charges)



```
plot(y=df_clean$sale_price,df_clean$sq_footage, xlab='Square Footage', ylab='Sales Price',)
```



```
library(ggplot2)
#plot(y=df_final$sale_price,df_clean$approx_year_built, xlab = 'Approx Year Built', ylab= 'Sales Price',
#geom_rug(sides = "bl")

#df_clean %<>%
#select(-sq_footage)

#Fix issue when using MLR, Notice that there is missingness so ill set it to other
df_clean$fuel_type[df_clean$fuel_type == 'none'] = 'other'
df_clean$fuel_type[is.na(df_clean$fuel_type)] = 'other'
df_clean$fuel_type = factor(df_clean$fuel_type)
#df_cleandd = df_clean[df_clean$sq_footage < 2500,]
#df_clean = subset(df_clean, df_clean$sq_footage <= 2500 & !is.na(df_clean$sq_footage))
#df_clean$num_full_bathrooms = as.factor(df_clean$num_full_bathrooms)
skim(df_clean)
```

Table 10: Data summary

Name	df_clean
Number of rows	528
Number of columns	16
Column type frequency:	
factor	5
numeric	11
Group variables	None

Variable type: factor

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
cats_allowed	0	1.00	FALSE	2	no: 285, yes: 243
coop_condo	0	1.00	FALSE	2	co-: 399, con: 129
dining_room_type	120	0.77	FALSE	4	com: 241, for: 116, oth: 49, din: 2
dogs_allowed	0	1.00	FALSE	2	no: 381, yes: 147
fuel_type	0	1.00	FALSE	4	gas: 301, oil: 180, oth: 36, ele: 11

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75
common_charges	396	0.25	433.92	205.40	70	288.50	390.5	537.7
garage_exists	0	1.00	0.18	0.38	0	0.00	0.0	0.0
maintenance_cost	142	0.73	821.85	378.77	155	639.25	734.0	880.0
approx_year_built	6	0.99	1962.38	20.56	1915	1950.00	1957.0	1968.0
num_bedrooms	0	1.00	1.54	0.75	0	1.00	1.0	2.0
num_floors_in_building	108	0.80	7.08	6.83	1	2.00	6.0	7.0
num_full_bathrooms	0	1.00	1.20	0.42	1	1.00	1.0	1.0
num_total_rooms	0	1.00	4.02	1.20	1	3.00	4.0	5.0
sq_footage	315	0.40	965.28	490.42	375	750.00	874.0	1010.0
sale_price	0	1.00	314956.56	179526.60	55000	171500.00	259500.0	428875.0
walk_score	0	1.00	83.10	13.09	15	76.00	85.0	94.0

```
set.seed(28)
M = tbl_df(apply(is.na(df_clean), 2, as.numeric))
colnames(M) = paste("is_missing_", colnames(df_clean), sep = "")
M %<>%
  select_if(function(x){sum(x) > 0})
head(M)

## # A tibble: 6 x 6
##   is_missing_comm~ is_missing_dini~ is_missing_main~ is_missing_appr~
##             <dbl>             <dbl>             <dbl>             <dbl>
## 1                0                0                1                0
## 2                1                0                0                0
## 3                0                0                1                0
## 4                0                0                1                0
## 5                1                0                0                0
## 6                1                0                0                0
## # ... with 2 more variables: is_missing_num_floors_in_building <dbl>,
## #   is_missing_sq_footage <dbl>


pacman::p_load(missForest)
dfimp = missForest(data.frame(df_clean))$ximp

## missForest iteration 1 in progress...done!
## missForest iteration 2 in progress...done!
## missForest iteration 3 in progress...done!
## missForest iteration 4 in progress...done!
## missForest iteration 5 in progress...done!
## missForest iteration 6 in progress...done!
## missForest iteration 7 in progress...done!
```


```
df_final = cbind(dfimp, M)
```

```
#skim(df_final)
```

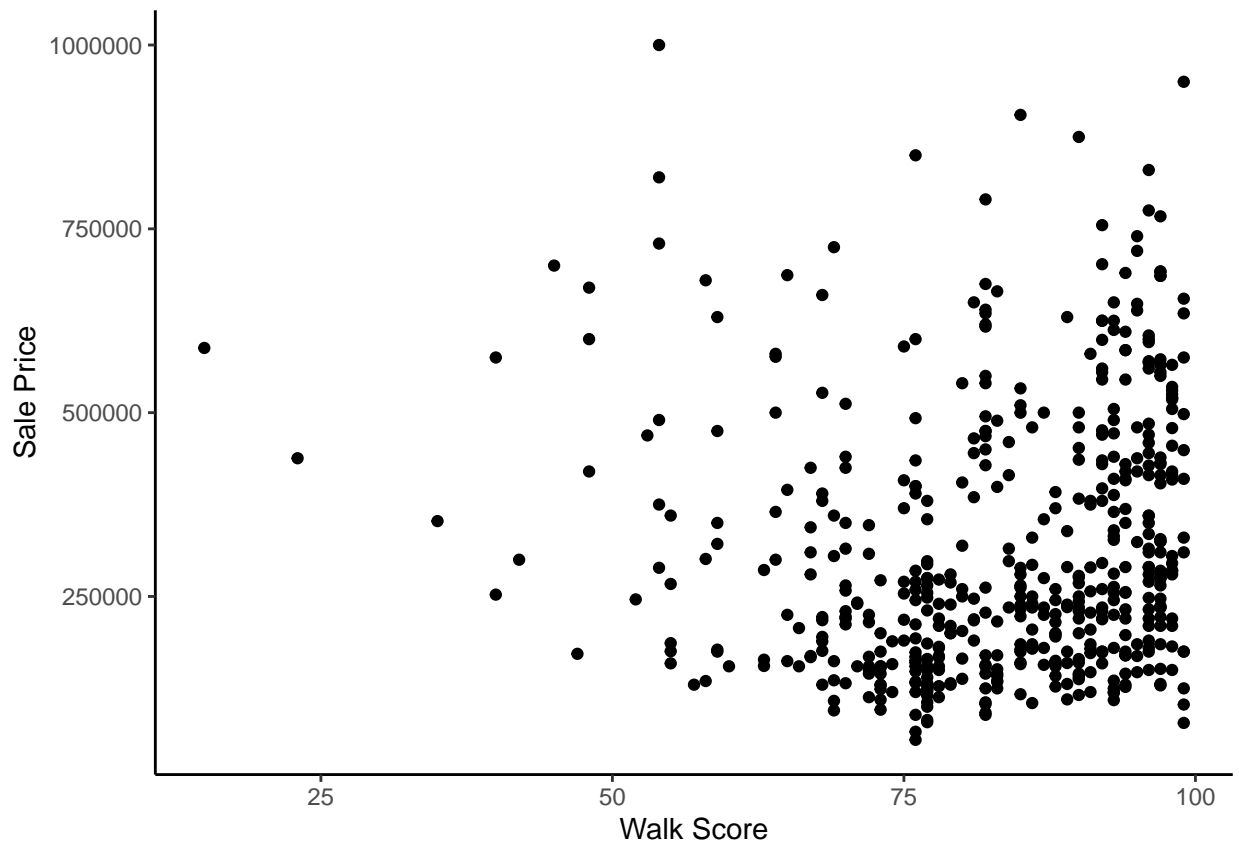
```
ggplot(df_final, aes(x=sq_footage, y=sale_price)) + geom_point() + theme_classic() +  
  labs(  
    x="Square Footage", y = "Sale Price")
```



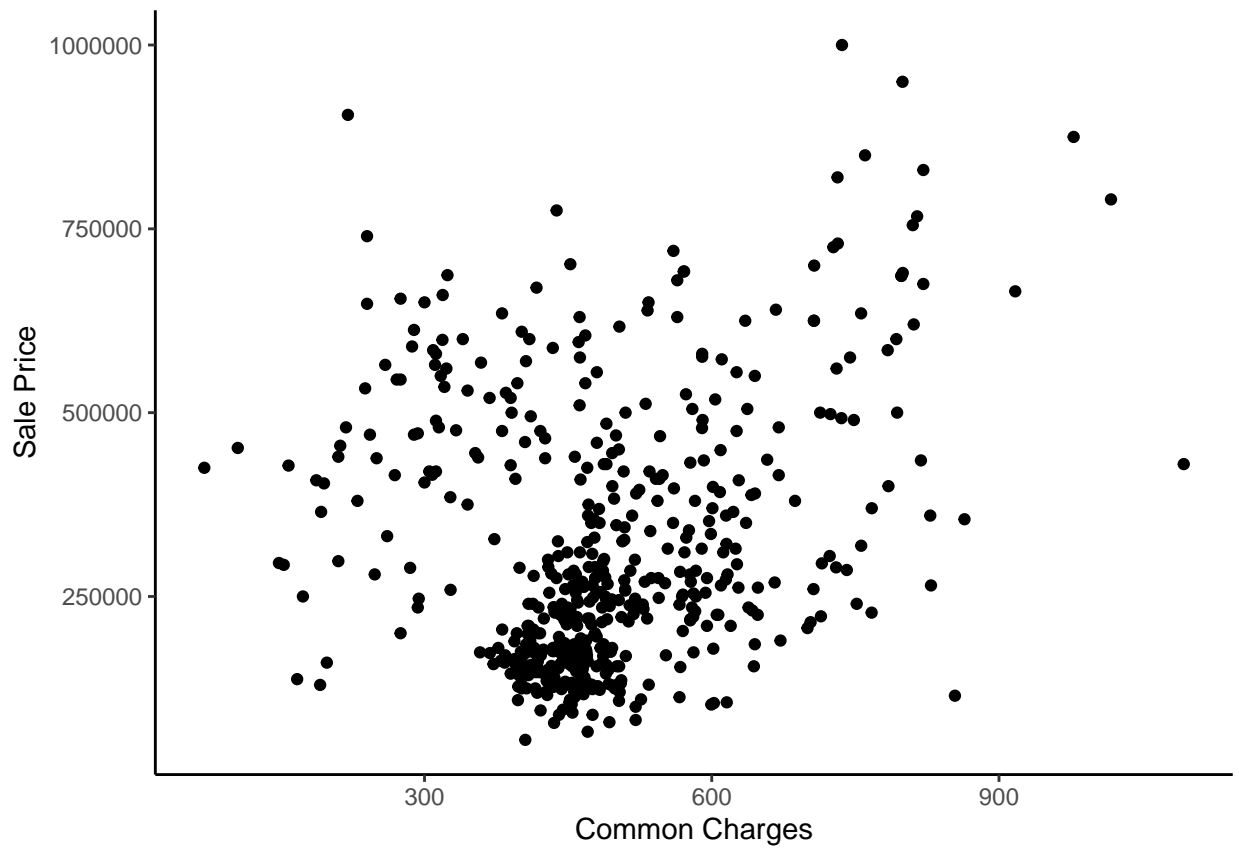
```
ggplot(df_final, aes(x=num_full_bathrooms, y=sale_price)) + geom_point() + theme_classic() +  
  labs(  
    x="Number Full Bathrooms", y = "Sale Price")
```



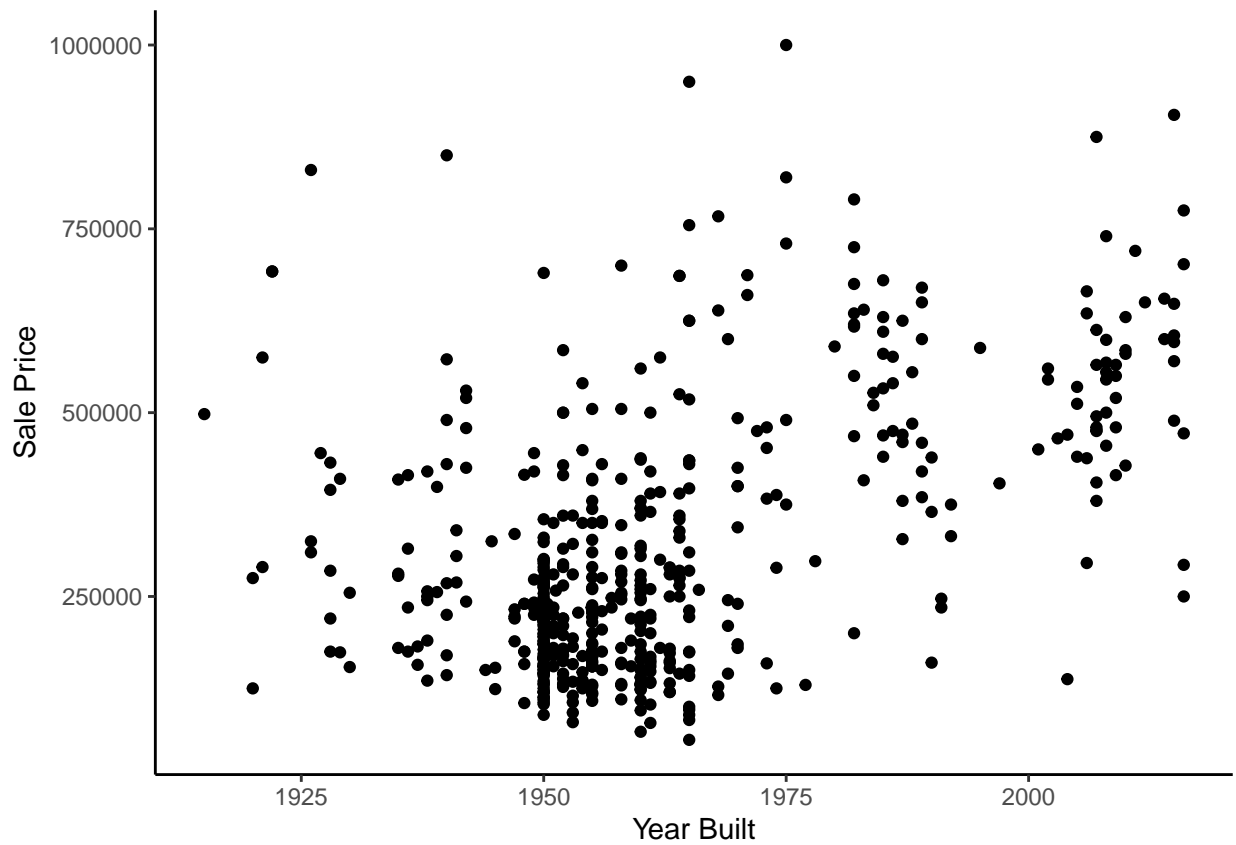
```
ggplot(df_final, aes(x=walk_score, y=sale_price)) + geom_point() + theme_classic() +  
  labs(  
    x="Walk Score", y = "Sale Price")
```



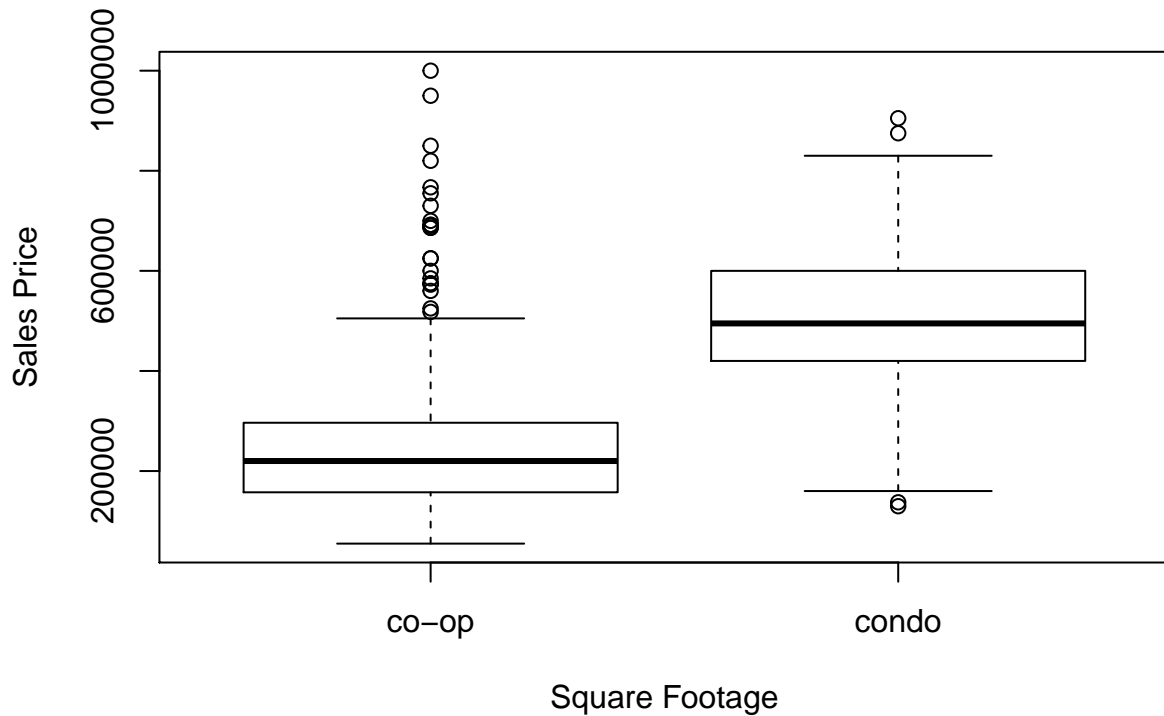
```
ggplot(df_final, aes(x=common_charges, y=sale_price)) + geom_point() + theme_classic() +  
  labs(  
    x="Common Charges", y = "Sale Price")
```

```
ggplot(df_final, aes(x=approx_year_built, y=sale_price)) + geom_point() + theme_classic() +  
  labs(  
    x="Year Built", y = "Sale Price")
```



```
plot(y=df_clean$sale_price,df_clean$coop_condo, xlab='Square Footage', ylab='Sales Price',)
```



```
# Tried to one hot encode data. MAJOR FAIL. R takes care of this since data is already factors
# df_dummy = copy(df_final)
```

```

# df_dummy$cats_allowed = model.matrix(~df_dummy$cats_allowed + 0)
# df_dummy$coop_condo = model.matrix(~df_dummy$coop_condo + 0)
# df_dummy$dining_room_type = model.matrix(~df_dummy$dining_room_type + 0)
# df_dummy$dogs_allowed = model.matrix(~df_dummy$dogs_allowed + 0)
# df_dummy$fuel_type = model.matrix(~df_dummy$fuel_type + 0)
# library(data.table, mltools)
# something = copy(df_final)
#
# something$fuel_type = cbind(model.matrix(~something$fuel_type))

```

```
colnames(df_final)
```

```

## [1] "cats_allowed"           "common_charges"
## [3] "coop_condo"             "dining_room_type"
## [5] "dogs_allowed"           "fuel_type"
## [7] "garage_exists"          "maintenance_cost"
## [9] "approx_year_built"      "num_bedrooms"
## [11] "num_floors_in_building" "num_full_bathrooms"
## [13] "num_total_rooms"        "sq_footage"
## [15] "sale_price"             "walk_score"
## [17] "is_missing_common_charges" "is_missing_dining_room_type"
## [19] "is_missing_maintenance_cost" "is_missing_approx_year_built"
## [21] "is_missing_num_floors_in_building" "is_missing_sq_footage"

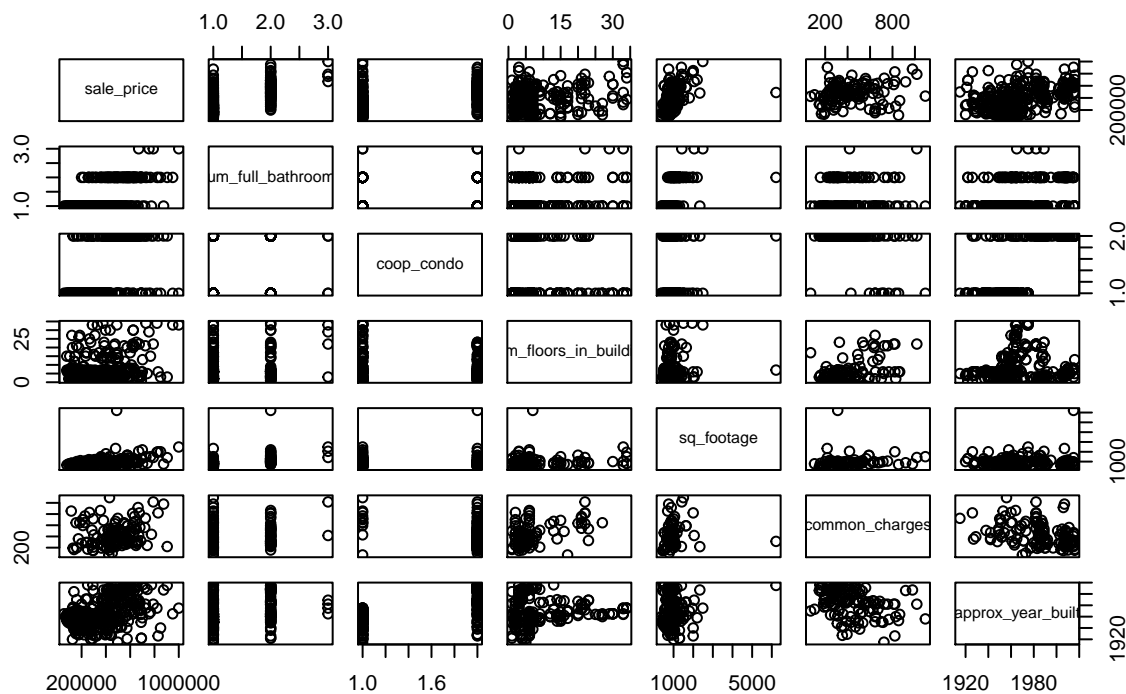
```

```

pairs(~sale_price+num_full_bathrooms+coop_condo+num_floors_in_building+sq_footage+common_charges+approx_year_built,
      main="Scatterplot Matrix")

```

Scatterplot Matrix



```
skim(df_final)
```

Table 13: Data summary

Name	df_final
Number of rows	528
Number of columns	22
Column type frequency:	
factor	5
numeric	17
Group variables	None

Variable type: factor

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
cats_allowed	0	1	FALSE	2	no: 285, yes: 243
coop_condo	0	1	FALSE	2	co-: 399, con: 129
dining_room_type	0	1	FALSE	4	com: 332, for: 139, oth: 55, din: 2
dogs_allowed	0	1	FALSE	2	no: 381, yes: 147
fuel_type	0	1	FALSE	4	gas: 301, oil: 180, oth: 36, ele: 11

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p
common_charges	0	1	489.07	138.73	70	417.43	469.
garage_exists	0	1	0.18	0.38	0	0.00	0.
maintenance_cost	0	1	810.71	361.55	155	605.50	722.
approx_year_built	0	1	1962.25	20.48	1915	1950.00	1956.
num_bedrooms	0	1	1.54	0.75	0	1.00	1.
num_floors_in_building	0	1	7.08	6.33	1	3.00	6.
num_full_bathrooms	0	1	1.20	0.42	1	1.00	1.
num_total_rooms	0	1	4.02	1.20	1	3.00	4.
sq_footage	0	1	901.76	364.21	375	722.87	835.
sale_price	0	1	314956.56	179526.60	55000	171500.00	259500.
walk_score	0	1	83.10	13.09	15	76.00	85.
is_missing_common_charges	0	1	0.75	0.43	0	0.75	1.
is_missing_dining_room_type	0	1	0.23	0.42	0	0.00	0.
is_missing_maintenance_cost	0	1	0.27	0.44	0	0.00	0.
is_missing_approx_year_built	0	1	0.01	0.11	0	0.00	0.
is_missing_num_floors_in_building	0	1	0.20	0.40	0	0.00	0.
is_missing_sq_footage	0	1	0.60	0.49	0	0.00	1.

```

set.seed(28)
train.control <- trainControl(method = "repeatedcv",
                              number = 10, repeats = 3)

prop_test = 0.10
test_indices = sample(1 : nrow(df_final), round((prop_test) * nrow(df_final)))
df_test = df_final[test_indices, ]
y_test = df_test$sale_price
X_test = cbind(1, df_test)

```

```

#X_test$sale_price = NULL

train_indices = setdiff(1 : nrow(df_final), test_indices)
df_train = df_final[train_indices, ]
y_train = df_train$sale_price
X_train = cbind(1, df_train)
#X_train$sale_price = NULL

n_train = nrow(X_train)

#mod = train(sale_price ~ ., df_final, trControl = train.control, method = "lm")
mod = lm(sale_price ~ ., df_final)
summary(mod)$r.squared

## [1] 0.77106

summary(mod)$sigma

## [1] 88012.24

summary(mod)

##
## Call:
## lm(formula = sale_price ~ ., data = df_final)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -349600  -49834   -2811   41151  354828
##
## Coefficients:
##              Estimate Std. Error t value
## (Intercept)    -155672.69   630147.20  -0.247
## cats_allowedyes      8259.94   10904.47   0.757
## common_charges       41.95     46.23   0.907
## coop_condocondo    269327.88  24446.08  11.017
## dining_room_typedining area    18414.34  63357.85   0.291
## dining_room_typeformal    13871.81  10362.41   1.339
## dining_room_typeother    23839.44  13955.36   1.708
## dogs_allowedyes    16391.72  12111.28   1.353
## fuel_typegas       5820.46  28113.12   0.207
## fuel_typeoil       4870.13  28854.99   0.169
## fuel_typeother    25405.43  31956.20   0.795
## garage_exists      2612.47  11269.62   0.232
## maintenance_cost     106.06     20.81   5.096
## approx_year_built     -45.63     321.37  -0.142
## num_bedrooms     57695.80   9432.52   6.117
## num_floors_in_building    5329.85    856.38   6.224
## num_full_bathrooms    74247.49  13549.78   5.480
## num_total_rooms       63.03    6248.70   0.010
## sq_footage         27.21     15.75   1.728
## walk_score         1299.14     316.12   4.110
## is_missing_common_charges  38818.06  24560.10   1.581
## is_missing_dining_room_type  8535.89   9581.08   0.891

```

```
## is_missing_maintenance_cost      -28263.28    22039.47   -1.282
## is_missing_approx_year_built      20074.99    37107.74    0.541
## is_missing_num_floors_in_building  19195.20    10071.58    1.906
## is_missing_sq_footage             -17027.67     8413.31   -2.024
##                                     Pr(>|t|)
## (Intercept)                       0.8050
## cats_allowedyes                    0.4491
## common_charges                     0.3646
## coop_condocondo                   < 0.0000000000000002 ***
## dining_room_typedining area        0.7714
## dining_room_typeformal             0.1813
## dining_room_typeother              0.0882 .
## dogs_allowedyes                    0.1765
## fuel_typegas                       0.8361
## fuel_typeoil                       0.8660
## fuel_typeother                     0.4270
## garage_exists                      0.8168
## maintenance_cost                   0.00000049094 ***
## approx_year_built                  0.8871
## num_bedrooms                       0.00000000192 ***
## num_floors_in_building              0.00000000103 ***
## num_full_bathrooms                 0.00000006756 ***
## num_total_rooms                    0.9920
## sq_footage                         0.0846 .
## walk_score                         0.00004627702 ***
## is_missing_common_charges          0.1146
## is_missing_dining_room_type        0.3734
## is_missing_maintenance_cost        0.2003
## is_missing_approx_year_built       0.5888
## is_missing_num_floors_in_building  0.0572 .
## is_missing_sq_footage              0.0435 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 88010 on 502 degrees of freedom
## Multiple R-squared:  0.7711, Adjusted R-squared:  0.7597
## F-statistic: 67.63 on 25 and 502 DF,  p-value: < 0.00000000000000022
```

```
set.seed(28)
```

```
mod =lm(sale_price ~., data.frame(df_train),set.seed(28))
```

```
summary(mod)$r.squared
```

```
## [1] 0.7649951
```

```
summary(mod)$sigma
```

```
## [1] 88712.26
```

```
y_hat = predict(mod,data.frame(X_test))
```

```
e = y_test - y_hat
```

```
Rsq_oos = (var(y_test) - var(e)) / var(y_test)
```

```
cat("My R Squared in sample is ",summary(mod)$r.squared, "My RSME is:", summary(mod)$sigma)
```

```
## My R Squared in sample is  0.7649951 My RSME is: 88712.26
```

```

cat("\nMy R Squared out of sample is ",Rsqr_oos, "My RSME is:", sd(e))

##
## My R Squared out of sample is 0.8046578 My RSME is: 85564.51
#plot(y_test,y_hat)
summary(mod)

##
## Call:
## lm(formula = sale_price ~ ., data = data.frame(df_train), subset = set.seed(28))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -353328  -46131   -1660   43358  353257
##
## Coefficients:
##              Estimate Std. Error t value
## (Intercept)    -298395.10   659717.40  -0.452
## cats_allowedyes      9080.70   11575.97   0.784
## common_charges       31.45    52.67   0.597
## coop_condocondo    268580.86   25419.10  10.566
## dining_room_typedining area    23700.90   89632.03   0.264
## dining_room_typeformal      7681.67   11135.26   0.690
## dining_room_typeother    32171.80   14691.74   2.190
## dogs_allowedyes    19982.52   12690.44   1.575
## fuel_typegas     -3832.21   29743.63  -0.129
## fuel_typeoil    -4827.69   30620.42  -0.158
## fuel_typeother    5305.87   33823.15   0.157
## garage_exists      802.17   12224.26   0.066
## maintenance_cost    121.29    26.87   4.514
## approx_year_built     33.66    336.42   0.100
## num_bedrooms    54268.57   10137.47   5.353
## num_floors_in_building    4939.70    918.37   5.379
## num_full_bathrooms    82887.84   14540.51   5.700
## num_total_rooms     378.61    6557.42   0.058
## sq_footage        23.79    16.32   1.458
## walk_score        1134.48    336.43   3.372
## is_missing_common_charges    42305.48   25406.07   1.665
## is_missing_dining_room_type    7221.66   10162.44   0.711
## is_missing_maintenance_cost  -26899.32   22966.64  -1.171
## is_missing_approx_year_built    23817.49   37540.65   0.634
## is_missing_num_floors_in_building    14964.93   10575.50   1.415
## is_missing_sq_footage    -15340.57    8925.11  -1.719
##
##              Pr(>|t|)
## (Intercept)    0.651266
## cats_allowedyes    0.433193
## common_charges    0.550712
## coop_condocondo    < 0.0000000000000002 ***
## dining_room_typedining area    0.791574
## dining_room_typeformal    0.490644
## dining_room_typeother    0.029053 *
## dogs_allowedyes    0.116050
## fuel_typegas    0.897541

```

```
## fuel_typeoil 0.874794
## fuel_typeother 0.875417
## garage_exists 0.947709
## maintenance_cost 0.0000081193 ***
## approx_year_built 0.920351
## num_bedrooms 0.0000001381 ***
## num_floors_in_building 0.0000001209 ***
## num_full_bathrooms 0.0000000217 ***
## num_total_rooms 0.953983
## sq_footage 0.145619
## walk_score 0.000811 ***
## is_missing_common_charges 0.096576 .
## is_missing_dining_room_type 0.477687
## is_missing_maintenance_cost 0.242126
## is_missing_approx_year_built 0.526113
## is_missing_num_floors_in_building 0.157745
## is_missing_sq_footage 0.086338 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 88710 on 449 degrees of freedom
## Multiple R-squared:  0.765, Adjusted R-squared:  0.7519
## F-statistic: 58.46 on 25 and 449 DF, p-value: < 0.000000000000000022

# pacman::p_load(ggplot2, mlr3, mlr)
# library(mlr3)
# library(mlr)
# modeling_task = makeRegrTask(data=data.frame(X_train), target='sale_price')
# algorithm = makeLearner("regr.lm")
# validation = makeResampleDesc("CV", iters = 5)
# #Having issues with fuel_type none
# res = resample(algorithm, modeling_task, validation, measures = list(rmse))
# res
# #average rsme somehow worse than above
# mean(res$measures.test$rmse)
```

REGRESSION TREES.

Here the trees overfit in sample but they did pretty decent out of sample but not better than OLS

```
options(java.parameters = "-Xmx4000m")

X_train_CART = X_train
X_train_CART$sale_price = NULL

X_test_CART = X_test
X_test_CART$sale_price = NULL
tree_model = YARFCART(X_train_CART, y_train, bootstrap_indices = 1 : n_train, calculate_oob_error = TRUE)

## YARF initializing with a fixed 1 trees...
## YARF factors created...
## YARF after data preprocessed... 31 total features...
## Beginning YARF regression model construction...done.
## Calculating OOB error...done.
```



```
illustrate_trees(tree_model, max_depth = 4, open_file = TRUE, margin_in_px=200, font_size=20, length_in_px=
get_tree_num_nodes_leaves_max_depths(tree_model)
```

```
## $num_nodes
## [1] 375
##
## $num_leaves
## [1] 188
##
## $max_depths
## [1] 25
```

```
#In Sample Error
```

```
y_hat_train = predict(tree_model, X_train)
```

```
## Warning in predict.YARF(tree_model, X_train): Prediction set column names did not match training set
## Attempting to subset to training set columns.
```

```
e_train = y_train - y_hat_train
rsme_train = sd(e_train)
rsquared_train = (var(y_train) - var(e_train)) / var(y_train)
```

```
#Out of Sample Error
```

```
y_hat_test = predict(tree_model, X_test)
```

```
## Warning in predict.YARF(tree_model, X_test): Prediction set column names did not match training set
## Attempting to subset to training set columns.
```

```
e_test = y_test - y_hat_test
rsme_test = sd(e_test)
rsquared_test = (var(y_test) - var(e_test)) / var(y_test)
```

```
cat("My R Squared in sample is ", rsquared_train, "My RSME is:", rsme_train)
```


```
## My R Squared in sample is 0.9882321 My RSME is: 19320.98
```

```
cat("\nMy R Squared out of sample is ", rsquared_test, "My RSME is:", rsme_test)
```

```
##
```

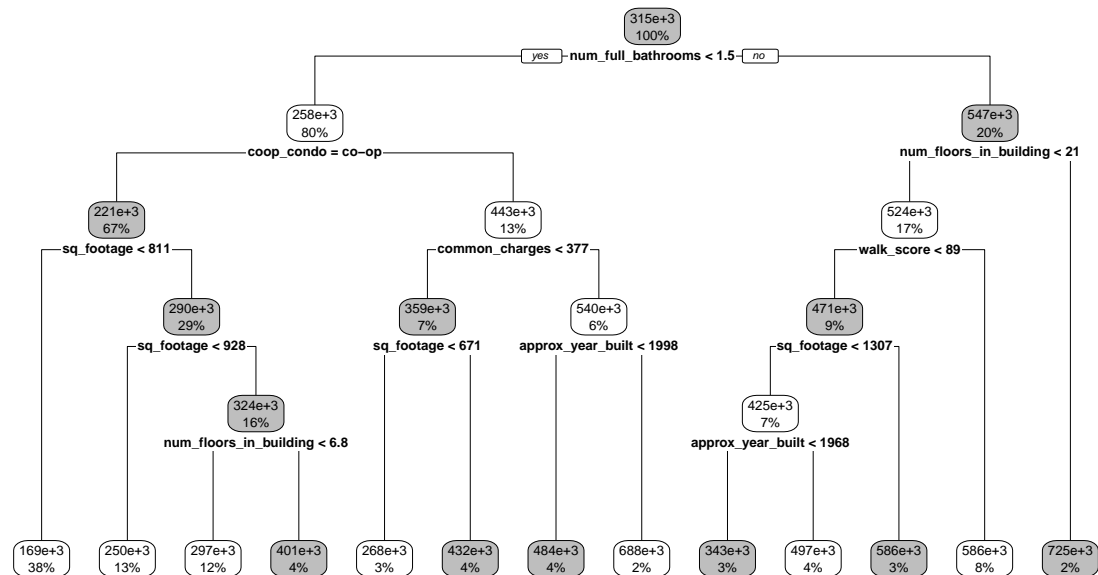
```
## My R Squared out of sample is 0.7961257 My RSME is: 87413.18
```

```
plot(y_test, y_hat)
```



```
library(rpart,mlr)
library("rpart.plot")
rpart_fit = rpart(sale_price ~., data.frame(X_train),method="anova")
rpart.plot(rpart_fit,box.col=c("grey", "white"))
```

```
## Warning: Bad 'data' field in model 'call' (expected a data.frame or a matrix).
## To silence this warning:
##   Call rpart.plot with roundint=FALSE,
##   or rebuild the rpart model with model=TRUE.
```



```
rpart_fit
```

```
## n= 475
##
## node), split, n, deviance, yval
```

```
##      * denotes terminal node
##
## 1) root 475 15036150000000 314877.8
##    2) num_full_bathrooms< 1.5 381 6589647000000 257531.4
##      4) coop_condo=co-op 318 2568349000000 220853.0
##        8) sq_footage< 811.1617 181 3894034000000 168632.4 *
##        9) sq_footage>=811.1617 137 1033253000000 289845.1
##          18) sq_footage< 928.3983 63 3366004000000 249738.1 *
##          19) sq_footage>=928.3983 74 5090368000000 323990.2
##            38) num_floors_in_building< 6.765 55 307402000000 297298.0 *
##            39) num_floors_in_building>=6.765 19 49014940000 401257.3 *
##      5) coop_condo=condo 63 1434082000000 442670.1
##        10) common_charges< 377 34 4946433000000 359439.4
##          20) sq_footage< 671.3365 15 92026640000 267975.9 *
##          21) sq_footage>=671.3365 19 178067100000 431647.4 *
##        11) common_charges>=377 29 4277701000000 540250.9
##          22) approx_year_built< 1997.5 21 1114547000000 483780.4 *
##          23) approx_year_built>=1997.5 8 735589900000 688486.0 *
##    3) num_full_bathrooms>=1.5 94 2115051000000 547313.8
##      6) num_floors_in_building< 20.5 83 1618691000000 523704.8
##        12) walk_score< 88.5 45 1045441000000 471400.0
##          24) sq_footage< 1307.457 32 4321335000000 425031.2
##            48) approx_year_built< 1967.5 15 98365230000 343366.7 *
##            49) approx_year_built>=1967.5 17 145464100000 497088.2 *
##          25) sq_footage>=1307.457 13 3751472000000 585538.5 *
##        13) walk_score>=88.5 38 3043495000000 585644.7 *
##      7) num_floors_in_building>=20.5 11 1010227000000 725454.5 *
```

```
library(mlr,mlr3)
```

```
## Loading required package: ParamHelpers
```

```
## 'mlr' is in maintenance mode since July 2019. Future development
## efforts will go into its successor 'mlr3' (<https://mlr3.ml-org.com>).
```

```
##
```

```
## Attaching package: 'mlr'
```

```
## The following object is masked from 'package:caret':
```

```
##
```

```
##      train
```

```
modeling_task = makeRegrTask(data=data.frame(X_train), target='sale_price')
```

```
## Warning in makeTask(type = type, data = data, weights = weights, blocking =
## blocking, : Empty factor levels were dropped for columns: dining_room_type
```

```
algorithm = makeLearner("regr.rpart")
```

```
validation = makeResampleDesc("CV", iters = 5)
```

```
#Having issues with fuel_type none
```

```
res = resample(algorithm, modeling_task, validation,measures = list(rmse))
```

```
## Resampling: cross-validation
```

```
## Measures:          rmse
```

```
## [Resample] iter 1: 77933.2398912
```

```
## [Resample] iter 2: 98589.0648772
```

```

## [Resample] iter 3:    99565.4579709
## [Resample] iter 4:    113289.5639943
## [Resample] iter 5:    106979.1754515
##
## Aggregated Result: rmse.test.rmse=99985.7422881
##
res

## Resample Result
## Task: data.frame(X_train)
## Learner: regr.rpart
## Aggr perf: rmse.test.rmse=99985.7422881
## Runtime: 0.069572
#average rsme somehow worse than above
mean(res$measures.test$rmse)

## [1] 99271.3
X_train_RF = X_train
X_train_RF$sale_price = NULL

X_test_RF = X_test
X_test_RF$sale_price = NULL
Bag_model = YARFBAG(X_train_RF, y_train, num_trees = 250, seed = 1 ,calculate_oob_error = TRUE)

## YARF initializing with a fixed 250 trees...
## YARF factors created...
## YARF after data preprocessed... 31 total features...
## Beginning YARF regression model construction...done.
## Calculating OOB error...done.
Bag_model$rmse_oob

## [1] 77425.65
y_hat_test_bag = predict(Bag_model, X_test)

## Warning in predict.YARF(Bag_model, X_test): Prediction set column names did not match training set c
## Attempting to subset to training set columns.
s_e_bag = sd(y_test - y_hat_test_bag)
s_e_bag

## [1] 76296.75
#In Sample Error
y_hat_train = predict(Bag_model,X_train)

## Warning in predict.YARF(Bag_model, X_train): Prediction set column names did not match training set c
## Attempting to subset to training set columns.
e_train = y_train - y_hat_train
rsme_train = sd(e_train)
rsquared_train = (var(y_train) - var(e_train)) / var(y_train)

```

```

#Out of Sample Error
y_hat_test = predict(Bag_model, X_test)

## Warning in predict.YARF(Bag_model, X_test): Prediction set column names did not match training set c
## Attempting to subset to training set columns.

e_test = y_test - y_hat_test
rsme_test = sd(e_test)
rsquared_test = (var(y_test) - var(e_test)) / var(y_test)

cat("My R Squared in sample is ",rsquared_train, "My RSME is:", rsme_train,"\n" )

## My R Squared in sample is  0.9702822 My RSME is: 30703.46
cat("\nMy R Squared out of sample is ",rsquared_test, "My RSME is:",rsme_test,"\n" )

##
## My R Squared out of sample is  0.8446824 My RSME is: 76296.75

df_final_bag_all = copy(df_final)
y_all = df_final_bag_all$sale_price
df_final_bag_all$sale_price = NULL
mod_bag_all = YARFBAG(df_final_bag_all, y_all, num_trees = 250, seed = 28)

## YARF initializing with a fixed 250 trees...
## YARF factors created...
## YARF after data preprocessed... 30 total features...
## Beginning YARF regression model construction...done.
## Calculating OOB error...done.

mod_bag_all$rmse_oob

## [1] 76421.52

```

RANDOM FOREST

```

set.seed(2)
X_train_RF = X_train
X_train_RF$sale_price = NULL

X_test_RF = X_test
X_test_RF$sale_price = NULL

RF_model = YARF(X_train_RF, y_train, num_trees = 250, seed = 1 ,calculate_oob_error = TRUE)

## YARF initializing with a fixed 250 trees...
## YARF factors created...
## YARF after data preprocessed... 31 total features...
## Beginning YARF regression model construction...done.
## Calculating OOB error...done.

#In Sample Error
y_hat_train = predict(RF_model,X_train)

## Warning in predict.YARF(RF_model, X_train): Prediction set column names did not match training set c
## Attempting to subset to training set columns.

```

```

e_train = y_train - y_hat_train
rsme_train = sd(e_train)
rsquared_train = (var(y_train) - var(e_train)) / var(y_train)

#Out of Sample Error
y_hat_test = predict(RF_model, X_test)

## Warning in predict.YARF(RF_model, X_test): Prediction set column names did not match training set columns
## Attempting to subset to training set columns.
e_test = y_test - y_hat_test
rsme_test = sd(e_test)
rsquared_test = (var(y_test) - var(e_test)) / var(y_test)

cat("My R Squared in sample is ",rsquared_train, "My RSME is:", rsme_train,"\n" )

## My R Squared in sample is 0.9657972 My RSME is: 32938.96
cat("\nMy R Squared out of sample is ",rsquared_test, "My RSME is:",rsme_test,"\n" )

##
## My R Squared out of sample is 0.820339 My RSME is: 82058.32
cat("OOB RSME:",RF_model$rmse_oob,"\n" )

## OOB RSME: 78439.13
cat("GAIN OVER TREES",(mod_bag_all$rmse_oob - RF_model$rmse_oob) / mod_bag_all$rmse_oob * 100, "%\n" )

## GAIN OVER TREES -2.640113 %
RF_model$rmse_oob

## [1] 78439.13
set.seed(28)
modeling_task = makeRegrTask(data=data.frame(X_train), target='sale_price')

## Warning in makeTask(type = type, data = data, weights = weights, blocking =
## blocking, : Empty factor levels were dropped for columns: dining_room_type
algorithm = makeLearner("regr.randomForest")
validation = makeResampleDesc("CV", iters = 5)
#Having issues with fuel_type none
res = resample(algorithm, modeling_task, validation,measures = list(rmse))

## Resampling: cross-validation
## Measures:          rmse
## [Resample] iter 1: 87545.1912787
## [Resample] iter 2: 75434.2169342
## [Resample] iter 3: 73396.4503821
## [Resample] iter 4: 82743.6328139
## [Resample] iter 5: 55317.9247182
##
## Aggregated Result: rmse.test.rmse=75694.2561831

```

```
##
res

## Resample Result
## Task: data.frame(X_train)
## Learner: regr.randomForest
## Aggr perf: rmse.test.rmse=75694.2561831
## Runtime: 2.80878

#average rsme somehow worse than above
mean(res$measures.test$rmse)

## [1] 74887.48

set.seed(28)
modeling_task = makeRegrTask(data=data.frame(X_train), target='sale_price')

## Warning in makeTask(type = type, data = data, weights = weights, blocking =
## blocking, : Empty factor levels were dropped for columns: dining_room_type

algorithm = makeLearner("regr.randomForest")
holdout = makeResampleDesc("Holdout")
#Having issues with fuel_type none
res = resample(algorithm, modeling_task, holdout,measures = list(rmse))

## Resampling: holdout
## Measures:          rmse
## [Resample] iter 1:  85520.2547615
##
## Aggregated Result: rmse.test.rmse=85520.2547615
##
res

## Resample Result
## Task: data.frame(X_train)
## Learner: regr.randomForest
## Aggr perf: rmse.test.rmse=85520.2547615
## Runtime: 0.40167

#average rsme somehow worse than above
mean(res$measures.test$rmse)

## [1] 85520.25

# library(rpart)
# library(rpart.plot)
# fit = rpart(sale_price ~., data.frame(X_train),method="anova")
# rpart.plot(fit)
# summary(fit)
# pred
#
# in_e = y_train - pred
# sd(in_e)
# (var(y_train) - var(e)) / var(y_train)
# e = y_test - pred
```

```

# sd(e)
#
# Rsq_oos = (var(y_test) - var(e)) / var(y_test)
# #sd(e)
# cat("My R Squared in sample is ",summary(mod)$r.squared, "My RSME is:", sd(in_e))
# cat("\nMy R Squared out of sample is ",Rsq_oos, "My RSME is:", sd(e))
# ...
# ```{r}
# library(randomForest)
# control <- trainControl(method="cv", number=10)
#
#
# RegressionTree1 = train(sale_price~., data=data.frame(X_train), method="rpart", trControl=control)
# y_hat = predict(object = RegressionTree1,newdata = data.frame(X_test))
# sqrt(mean((y_hat-y_test)^2))
#
# RegressionTree = train(sale_price~., data=df_final, method="rpart", trControl=control)
# print(RegressionTree)
#
#
# ##
#
# fit = rpart(sale_price ~., data.frame(X_train),method = 'anova')
# printcp(fit)
# rpart.plot(fit)
# summary(fit)
# y_hat = predict(object = fit,newdata = data.frame(X_test))
# sqrt(mean((y_hat-y_test)^2))

# RandomForest = train(sale_price~., data=df_final, method="rf", trControl=control)
# print(RandomForest)
#
#

```