

# Lab 5

Tyron Samaroo

11:59PM March 7, 2020

Load the Boston housing data frame and create the vector  $y$  (the median value) and matrix  $X$  (all other features) from the data frame. Name the columns the same as Boston except for the first name it “(Intercept)”.

```
y = MASS::Boston$medv
X = MASS::Boston[, 1:13]
```

Run the OLS linear model to get  $b$ , the vector of coefficients. Do not use `lm`. This is  $(XX^T)^{-1}X^Ty$

```
X = cbind(1, as.matrix(X))
b = solve(t(X) %*% X) %*% t(X) %*% y
```

Find the hat matrix for this regression  $H$ . Verify its dimension is correct and verify its rank is correct.  $H = X(X^TX)^{-1}X^T$

```
H = X %*% solve(t(X) %*% X) %*% t(X)
dim(H)
```

```
## [1] 506 506
```

```
pacman::p_load(Matrix)
rankMatrix(H)
```

```
## [1] 14
## attr(,"method")
## [1] "tolNorm2"
## attr(,"useGrad")
## [1] FALSE
## attr(,"tol")
## [1] 1.123546e-13
```

Verify this is a projection matrix by verifying the two sufficient conditions. Use the `testthat` library's `expect_equal(matrix1, matrix2, tolerance = 1e-2)`.

```
pacman::p_load(testthat)
expect_equal(H, t(H))
expect_equal(H, H %*% H)
```

Find the matrix that projects onto the space of residuals  $H_{comp}$  and find its rank. Is this rank expected?

```
I = diag(nrow(H))
Hcomp = I - H
rankMatrix(Hcomp)
```

```
## [1] 14
## attr(,"method")
## [1] "tolNorm2"
## attr(,"useGrad")
## [1] FALSE
```

```
## attr("tol")
## [1] 1.123546e-13
rankMatrix(Hcomp, tol=1e-2)
```

```
## [1] 492
## attr("method")
## [1] "tolNorm2"
## attr("useGrad")
## [1] FALSE
## attr("tol")
## [1] 0.01
```

Verify this is a projection matrix by verifying the two sufficient conditions. Use the `testthat` library.

```
expect_equal(Hcomp, t(Hcomp))
expect_equal(Hcomp, Hcomp %*% Hcomp)
```

Use `diag` to find the trace of both `H` and `Hcomp`.

```
sum(diag(H))
```

```
## [1] 14
```

```
sum(diag(Hcomp))
```

```
## [1] 492
```

Do you have a conjecture about the trace of an orthogonal projection matrix?

trace is equal to the rank

Find the eigendecomposition of both `H` and `Hcomp` as `eigenvals_H`, `eigenvecs_H`, `eigenvals_Hcomp`, `eigenvecs_Hcomp`. Verify these results are the correct dimensions.

```
eigen_H = eigen(H)
eigen_Hcomp = eigen(Hcomp)

eigenvals_H = eigen_H$values
eigenvecs_H = eigen_H$vectors
eigenvals_Hcomp = eigen_Hcomp$values
eigenvecs_Hcomp = eigen_Hcomp$vectors
```

```
length(eigenvals_H)
```

```
## [1] 506
```

```
dim(eigenvecs_H)
```

```
## [1] 506 506
```

```
length(eigenvals_Hcomp)
```

```
## [1] 506
```

```
dim(eigenvecs_Hcomp)
```

```
## [1] 506 506
```

The eigendecomposition suffers from numerical error which is making them become imaginary. We can coerce imaginary numbers back to real by using the `Re` function. There is also lots of numerical error. Use the `Re` function to coerce to real and the `round` function to round all four objects to the nearest 10 digits.

Print out the eigenvalues of both  $H$  and  $H_{\text{comp}}$ . Is this expected?

[illegible][illegible]

```
apply(eigenvecs_H, MARGIN = 2, FUN = function(v){
  sqrt(sum(v^2))
})
```

3

```

## [71] 0.7375631 0.7200799 0.7200799 0.7067433 0.7067433 1.0000000 0.7291502
## [78] 0.7291502 0.9861821 0.9861821 0.7181425 0.7181425 0.7289067 0.7289067
## [85] 0.7069350 0.7069350 0.7147271 0.7147271 0.7278385 0.7278385 0.7016417
## [92] 0.7016417 0.6985773 0.6985773 0.6978982 0.6978982 0.6983338 0.6983338
## [99] 0.7001943 0.7001943 0.7051763 0.7051763 0.7555366 0.7555366 1.0000000
## [106] 0.6876725 0.6876725 0.7150475 0.7150475 0.7486783 0.7486783 0.7261142
## [113] 0.7261142 0.6985145 0.6985145 0.6836561 0.6836561 1.0000000 0.7187066
## [120] 0.7187066 0.6905980 0.6905980 0.7179283 0.7179283 0.7222461 0.7222461
## [127] 0.7207965 0.7207965 0.7004181 0.7004181 0.7014897 0.7014897 0.6990003
## [134] 0.6990003 0.7115622 0.7115622 0.7237343 0.7237343 0.7703283 0.7703283
## [141] 0.7899225 0.7899225 0.7259246 0.7259246 0.7076892 0.7076892 0.7727441
## [148] 0.7727441 0.7179819 0.7179819 0.7213128 0.7213128 0.6696315 0.6696315
## [155] 0.7021261 0.7021261 0.7045330 0.7045330 0.7363510 0.7363510 1.0000000
## [162] 0.6941660 0.6941660 0.7753994 0.7753994 0.7540773 0.7540773 0.7070886
## [169] 0.7070886 0.6903144 0.6903144 0.6988027 0.6988027 0.7120509 0.7120509
## [176] 0.7158204 0.7158204 0.7312673 0.7312673 1.0000000 0.6920332 0.6920332
## [183] 0.6866946 0.6866946 0.7373332 0.7373332 0.7270449 0.7270449 0.7073476
## [190] 0.7073476 0.7085796 0.7085796 0.7045089 0.7045089 0.7563864 0.7563864
## [197] 0.7618808 0.7618808 0.7044355 0.7044355 0.8477404 0.8477404 0.7109985
## [204] 0.7109985 0.6987366 0.6987366 0.7708438 0.7708438 0.6986430 0.6986430
## [211] 0.7336177 0.7336177 1.0000000 0.6825439 0.6825439 0.6940425 0.6940425
## [218] 0.7314948 0.7314948 0.7171839 0.7171839 0.7281819 0.7281819 1.0000000
## [225] 0.7031402 0.7031402 0.6951803 0.6951803 0.7152899 0.7152899 0.7281453
## [232] 0.7281453 0.7160614 0.7160614 0.7208316 0.7208316 0.7089619 0.7089619
## [239] 0.7120048 0.7120048 0.7095858 0.7095858 0.8786539 0.8786539 0.6955125
## [246] 0.6955125 0.7157227 0.7157227 1.0000000 0.7616379 0.7616379 0.7383818
## [253] 0.7383818 0.7198086 0.7198086 0.7054253 0.7054253 0.7185741 0.7185741
## [260] 0.7113626 0.7113626 0.7210185 0.7210185 0.7061893 0.7061893 0.9134175
## [267] 0.9134175 0.7129986 0.7129986 0.7232002 0.7232002 0.7393014 0.7393014
## [274] 0.7209286 0.7209286 0.7185819 0.7185819 0.7586169 0.7586169 0.7560925
## [281] 0.7560925 0.7144534 0.7144534 0.7047713 0.7047713 0.7334610 0.7334610
## [288] 0.7207679 0.7207679 0.6943252 0.6943252 1.0000000 0.7289015 0.7289015
## [295] 0.7186150 0.7186150 0.7316530 0.7316530 0.8856238 0.8856238 0.7071873
## [302] 0.7071873 0.7283910 0.7283910 0.7209085 0.7209085 0.7375309 0.7375309
## [309] 0.7114813 0.7114813 0.7944529 0.7944529 0.7303951 0.7303951 0.7186130
## [316] 0.7186130 0.8186602 0.8186602 0.6977128 0.6977128 0.7075066 0.7075066
## [323] 0.7688476 0.7688476 0.6874561 0.6874561 0.7118374 0.7118374 1.0000000
## [330] 0.6779750 0.6779750 0.7146818 0.7146818 0.7260265 0.7260265 0.7126152
## [337] 0.7126152 1.0000000 0.7150297 0.7150297 0.7228253 0.7228253 0.7251544
## [344] 0.7251544 0.7112836 0.7112836 0.7131162 0.7131162 0.7254483 0.7254483
## [351] 0.7059673 0.7059673 0.7087931 0.7087931 0.7081581 0.7081581 0.7231071
## [358] 0.7231071 0.7004617 0.7004617 0.8706722 0.8706722 0.7317309 0.7317309
## [365] 0.7945923 0.7945923 1.0000000 0.7322846 0.7322846 0.7326250 0.7326250
## [372] 0.7174917 0.7174917 0.7517336 0.7517336 0.7375921 0.7375921 0.6711866
## [379] 0.6711866 0.6895559 0.6895559 0.6880716 0.6880716 0.7022134 0.7022134
## [386] 0.7630476 0.7630476 0.7072201 0.7072201 0.7656041 0.7656041 0.6909803
## [393] 0.6909803 0.6955879 0.6955879 0.7054902 0.7054902 0.8115084 0.8115084
## [400] 0.6961595 0.6961595 1.0000000 0.6879843 0.6879843 0.7306674 0.7306674
## [407] 0.8083010 0.8083010 0.8121622 0.8121622 0.6934602 0.6934602 0.7164527
## [414] 0.7164527 0.7074909 0.7074909 0.7179766 0.7179766 0.7076450 0.7076450
## [421] 0.7497639 0.7497639 1.0000000 0.8140732 0.8140732 0.7093069 0.7093069
## [428] 0.9117919 0.9117919 0.9473980 0.9473980 0.7074843 0.7074843 0.7350700
## [435] 0.7350700 0.7417871 0.7417871 0.7547830 0.7547830 0.7366121 0.7366121
## [442] 0.7769113 0.7769113 0.7810845 0.7810845 0.6892008 0.6892008 0.8055246

```

```
## [449] 0.8055246 0.6814249 0.6814249 0.7272492 0.7272492 0.7020231 0.7020231
## [456] 0.7611428 0.7611428 0.7367486 0.7367486 0.7048714 0.7048714 0.7400683
## [463] 0.7400683 1.0000000 0.7952697 0.7952697 0.6912670 0.6912670 0.8435088
## [470] 0.8435088 0.7646925 0.7646925 0.7883719 0.7883719 0.8329977 0.8329977
## [477] 0.7629826 0.7629826 1.0000000 0.7070124 0.7070124 0.7649891 0.7649891
## [484] 0.6532954 0.6532954 0.7107735 0.7107735 1.0000000 0.7442381 0.7442381
## [491] 1.0000000 0.7999522 0.7999522 1.0000000 0.7369916 0.7369916 0.7339597
## [498] 0.7339597 0.9876976 0.9876976 0.9150144 0.9150144 0.8738908 0.8738908
## [505] 1.0000000 1.0000000
```

Is this expected? What is the convention for eigenvectors in R's `eigen` function?

Yes. The convention is length 1.

The first  $p+1$  eigenvectors are the columns of  $X$  but they are in arbitrary order. Find the column that represents the one-vector.

```
head(eigenvecs_H[, 3])
```

```
## [1] 0.04561114 0.04803979 0.04756035 0.04872081 0.04788055 0.04985263
```

Why is it not exactly 506 1's?

Numeric error

Use the first  $p+1$  eigenvectors as a model matrix and run the OLS model of `medv` on that model matrix.

```
mod1 = lm(y ~ X)
mod2 = lm(y ~ eigenvecs_H[, 1:14])
summary(mod1)
```

```
##
## Call:
## lm(formula = y ~ X)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.595  -2.730  -0.518   1.777   26.199
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.646e+01  5.103e+00   7.144 3.28e-12 ***
## X              NA          NA      NA      NA
## Xcrim        -1.080e-01  3.286e-02  -3.287 0.001087 **
## Xzn           4.642e-02  1.373e-02   3.382 0.000778 ***
## Xindus        2.056e-02  6.150e-02   0.334 0.738288
## Xchas         2.687e+00  8.616e-01   3.118 0.001925 **
## Xnox        -1.777e+01  3.820e+00  -4.651 4.25e-06 ***
## Xrm           3.810e+00  4.179e-01   9.116 < 2e-16 ***
## Xage          6.922e-04  1.321e-02   0.052 0.958229
## Xdis        -1.476e+00  1.995e-01  -7.398 6.01e-13 ***
## Xrad          3.060e-01  6.635e-02   4.613 5.07e-06 ***
## Xtax        -1.233e-02  3.760e-03  -3.280 0.001112 **
## Xptratio     -9.527e-01  1.308e-01  -7.283 1.31e-12 ***
## Xblack        9.312e-03  2.686e-03   3.467 0.000573 ***
## Xlstat       -5.248e-01  5.072e-02 -10.347 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 4.745 on 492 degrees of freedom
## Multiple R-squared:  0.7406, Adjusted R-squared:  0.7338
## F-statistic: 108.1 on 13 and 492 DF,  p-value: < 2.2e-16
```

```
summary(mod2)
```

```
##
## Call:
## lm(formula = y ~ eigenvecs_H[, 1:14])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.5885  -2.7528  -0.5003   1.6738  26.1135
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.874e+08  3.403e+08   0.845   0.399
## eigenvecs_H[, 1:14]1  1.238e+08  1.466e+08   0.845   0.399
## eigenvecs_H[, 1:14]2 -2.909e+08  3.444e+08  -0.845   0.399
## eigenvecs_H[, 1:14]3 -6.559e+09  7.766e+09  -0.845   0.399
## eigenvecs_H[, 1:14]4  6.044e+08  7.156e+08   0.845   0.399
## eigenvecs_H[, 1:14]5  9.666e+07  1.144e+08   0.845   0.399
## eigenvecs_H[, 1:14]6 -3.000e+08  3.551e+08  -0.845   0.399
## eigenvecs_H[, 1:14]7  2.716e+08  3.216e+08   0.845   0.399
## eigenvecs_H[, 1:14]8  1.631e+07  1.931e+07   0.845   0.399
## eigenvecs_H[, 1:14]9 -9.690e+08  1.147e+09  -0.845   0.399
## eigenvecs_H[, 1:14]10 -8.130e+06  9.626e+06  -0.845   0.399
## eigenvecs_H[, 1:14]11 -2.030e+08  2.404e+08  -0.845   0.399
## eigenvecs_H[, 1:14]12  1.552e+08  1.838e+08   0.845   0.399
## eigenvecs_H[, 1:14]13  3.570e+08  4.227e+08   0.845   0.399
## eigenvecs_H[, 1:14]14 -3.189e+07  3.776e+07  -0.845   0.399
##
## Residual standard error: 4.747 on 491 degrees of freedom
## Multiple R-squared:  0.741, Adjusted R-squared:  0.7336
## F-statistic: 100.3 on 14 and 491 DF,  p-value: < 2.2e-16
```

Is  $b$  about the same above (in arbitrary order)?

NO, the eigen vectors are scaled to be unit length

Calculate  $\hat{y}$  using the hat matrix.

```
y_hat= H %*% y
```

Calculate  $e$  two ways: (1) the difference of  $y$  and  $\hat{y}$  and (2) the projection onto the space of the residuals. Verify the two means of calculating the residuals provide the same results via `expect_equal`.

```
e1 = y - y_hat
e2 = Hcomp %*% y
expect_equal(e1, e2)
```

Calculate  $R^2$  using the angle relationship between the responses and their predictions.  $\$$

```
length_of_vec = function(v){sqrt(sum(v^2))}
y_avg_adj = y - mean(y)
y_yhat_adj = y_hat - mean(y)
(sum(y_avg_adj *
```

```

    y_yhat_adj
  )
  /
  (length_of_vec(y_avg_adj) *
    length_of_vec(y_yhat_adj))
  ) ** 2

```

```
## [1] 0.7406427
```

```

(
  sum(
    y_avg_adj * y_yhat_adj
  )
  /
  (
    length_of_vec(y_avg_adj) * length_of_vec(y_yhat_adj)) ^ 2

```

```
## [1] 0.7406427
```

Find the cosine-squared of  $y - \bar{y}$  and  $\hat{y} - \bar{y}$  and verify it is the same as  $R^2$ .

```
summary(mod1)$r.squared
```

```
## [1] 0.7406427
```

```

theta_in_rad = cos(
  ((y_avg_adj) %*% y_yhat_adj)
  /
  (length_of_vec(y_avg_adj) * length_of_vec(y_yhat_adj))
)
(theta_in_rad * 180 / pi)

```

```
## [1] 37.35559
```

```
theta_in_rad ^ 2
```

```
## [1] 0.4250755
```

```
## [1,] 0.4250755
```

Verify  $\hat{y}$  and  $e$  are orthogonal.

```
sum(y_hat*e1)
```

```
## [1] -4.991219e-08
```

Verify  $\hat{y} - \bar{y}$  and  $e$  are orthogonal.

```
sum((y_hat - mean(y)) * e1)
```

```
## [1] 2.832455e-09
```

Verify the sum of squares identity which we learned was due to the Pythagorean Theorem (applies since the projection is specifically orthogonal). You need to compute all three quantities first: SST, SSR and SSE.

```

SSE = sum((y - y_hat)^2)
SST = sum((y - mean(y))^2)
SSR = sum((y_hat - mean(y))^2)
1 - (SSE / SST)

```

```
## [1] 0.7406427
```

```
expect_equal(SSR/SST, 1 - (SSE / SST), (SST-SSE) / SST)
a = sqrt(SSR)
b = sqrt(SSE)
c = sqrt(SST)
expect_equal(c^2, a^2 + b^2)
```

Create a matrix that is  $(p + 1) \times (p + 1)$  full of NA's. Label the columns the same columns as X. Do not label the rows. For the first row, find the OLS estimate of the  $y$  regressed on the first column only and put that in the first entry. For the second row, find the OLS estimates of the  $y$  regressed on the first and second columns of  $X$  only and put them in the first and second entries. For the third row, find the OLS estimates of the  $y$  regressed on the first, second and third columns of  $X$  only and put them in the first, second and third entries, etc. For the last row, fill it with the full OLS estimates.

```
new_matrix= matrix(NA,ncol(X),ncol(X))
dim(new_matrix)
```

```
## [1] 14 14
```

```
colnames(new_matrix) = colnames(X)
```

```
# for(i in 1:ncol(new_matrix)){
#   matrix_col_i = X[, 1:i, drop = FALSE] # entire column
#   new_matrix[i, 1:i] = solve(t(matrix_col_i) %*% matrix_col_i ) %*% t(matrix_col_i) %*% y #regression
# }
```

```
for(i in 1:ncol(new_matrix)){
  matrix_col_i = X[, 1:i, drop = FALSE]
  for(j in 1:i){
    b = solve(t(matrix_col_i) %*% matrix_col_i ) %*% t(matrix_col_i) %*% y
    new_matrix[i, 1:i] = b
  }
}
```

```
#new_matrix
```

```
round(new_matrix,2) # easier to view
```

```
##          crim    zn indus chas      nox    rm    age    dis    rad    tax ptratio
## [1,]  22.53    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA
## [2,]  24.03 -0.42    NA    NA    NA    NA    NA    NA    NA    NA    NA
## [3,]  22.49 -0.35  0.12    NA    NA    NA    NA    NA    NA    NA    NA
## [4,]  27.39 -0.25  0.06 -0.42    NA    NA    NA    NA    NA    NA    NA
## [5,]  27.11 -0.23  0.06 -0.44  6.89    NA    NA    NA    NA    NA    NA
## [6,]  29.49 -0.22  0.06 -0.38  7.03 -5.42    NA    NA    NA    NA    NA
## [7,] -17.95 -0.18  0.02 -0.14  4.78 -7.18  7.34    NA    NA    NA    NA
## [8,] -18.26 -0.17  0.01 -0.13  4.84 -4.36  7.39 -0.02    NA    NA    NA
## [9,]   0.83 -0.20  0.06 -0.23  4.58 -14.45  6.75 -0.06 -1.76    NA    NA
## [10,]  0.16 -0.18  0.06 -0.21  4.54 -13.34  6.79 -0.06 -1.75 -0.05    NA
## [11,]  2.99 -0.18  0.07 -0.10  4.11 -12.59  6.66 -0.05 -1.73  0.16 -0.01    NA
## [12,] 27.15 -0.18  0.04 -0.04  3.49 -22.18  6.08 -0.05 -1.58  0.25 -0.01 -1.00
## [13,] 20.65 -0.16  0.04 -0.03  3.22 -20.48  6.12 -0.05 -1.55  0.28 -0.01 -1.01
## [14,] 36.46 -0.11  0.05  0.02  2.69 -17.77  3.81  0.00 -1.48  0.31 -0.01 -0.95
##      black lstat
## [1,]    NA    NA
## [2,]    NA    NA
## [3,]    NA    NA
## [4,]    NA    NA
```



```
## [5,]    NA    NA
## [6,]    NA    NA
## [7,]    NA    NA
## [8,]    NA    NA
## [9,]    NA    NA
## [10,]   NA    NA
## [11,]   NA    NA
## [12,]   NA    NA
## [13,]  0.01    NA
## [14,]  0.01 -0.52
```

Examine this matrix. Why are the estimates changing from row to row as you add in more predictors

We are examining one feature at a time and seeing how far its away from the mean.

Clear the workspace and load the diamonds dataset in the package `ggplot2`.

```
rm(list=ls())
pacman::p_load(ggplot2)
data("diamonds")
```

Extract `y`, the price variable and `col`, the nominal variable “color” as vectors.

```
y = diamonds$price
col = diamonds$color
```

Convert the `col` vector to `X` which contains an intercept and an appropriate number of dummies. Let the color G be the reference category as it is the modal color. Name the columns of `X` appropriately. The first should be “(Intercept)”. Delete `col`.

```
#Problem col is ordered factor cant just use relevel
# col = droplevels(col,"G")
# X = cbind(1,col)
#
# X
X = matrix(1,nrow(diamonds)) #

for(lev in levels(col)){
  if(lev != "G"){
    X = cbind(X, col == lev)
  }
}
colnames(X) = c("Intercept","D","E","F","H","I","J")
```

Repeat the iterative exercise above we did for Boston here.

```
new_matrix= matrix(NA,ncol(X),ncol(X))
colnames(new_matrix) = colnames(X)
for(i in 1:ncol(new_matrix)){
  matrix_col_i = X[, 1:i, drop = FALSE]
  for(j in 1:i){
    b = solve(t(matrix_col_i) %*% matrix_col_i ) %*% t(matrix_col_i) %*% y
    new_matrix[i, 1:i] = b
  }
}

price_model = lm(price ~ color, diamonds)
```

```
round(new_matrix,2)
```

```
##      Intercept      D      E      F      H      I      J
## [1,] 3932.80      NA      NA      NA      NA      NA      NA
## [2,] 4042.38 -872.42      NA      NA      NA      NA      NA
## [3,] 4295.54 -1125.59 -1218.79      NA      NA      NA      NA
## [4,] 4491.23 -1321.28 -1414.48 -766.34      NA      NA      NA
## [5,] 4493.17 -1323.22 -1416.42 -768.28 -6.50      NA      NA
## [6,] 4262.94 -1092.99 -1186.19 -538.06 223.72 828.93      NA
## [7,] 3999.14 -829.18 -922.38 -274.25 487.53 1092.74 1324.68
```

```
summary(price_model)
```

```
##
## Call:
## lm(formula = price ~ color, data = diamonds)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4989   -2619   -1376    1374   15654
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  4124.73      18.64  221.294 < 2e-16 ***
## color.L      2126.73      57.02   37.295 < 2e-16 ***
## color.Q       200.50      54.26    3.695 0.00022 ***
## color.C     -254.36      51.08   -4.979 6.41e-07 ***
## color^4        40.88      46.92    0.871 0.38361
## color^5     -228.88      44.36   -5.160 2.48e-07 ***
## color^6        87.92      40.22    2.186 0.02880 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3927 on 53933 degrees of freedom
## Multiple R-squared:  0.03128,    Adjusted R-squared:  0.03117
## F-statistic: 290.2 on 6 and 53933 DF,  p-value: < 2.2e-16
```

Why didn't the estimates change as we added more and more features?

The estimates did change we we added more featues.

Model price with both color and clarity with and without an intercept and report the coefficients.

```
with_inter_diamond_price_model = lm(price ~ color + clarity, diamonds)
```

```
without_inter_diamond_price_model = lm(price ~ 0 +color + clarity, diamonds)
summary(with_inter_diamond_price_model)
```

```
##
## Call:
## lm(formula = price ~ color + clarity, data = diamonds)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -6046   -2469   -1308    1164   16858
##
## Coefficients:
```

```
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) 3813.01      26.89 141.775 < 2e-16 ***
## color.L     2137.54      56.34  37.941 < 2e-16 ***
## color.Q      16.27      53.83   0.302 0.762406
## color.C     -221.70      50.39  -4.400 1.09e-05 ***
## color^4      77.36      46.28   1.671 0.094668 .
## color^5     -259.95      43.73  -5.945 2.78e-09 ***
## color^6       2.12      39.76   0.053 0.957479
## clarity.L   -1718.62     97.41 -17.642 < 2e-16 ***
## clarity.Q    -594.15     95.30  -6.235 4.56e-10 ***
## clarity.C     681.36     81.97   8.313 < 2e-16 ***
## clarity^4    -248.60     65.71  -3.783 0.000155 ***
## clarity^5     806.88     53.76  15.010 < 2e-16 ***
## clarity^6    -228.71     46.90  -4.876 1.08e-06 ***
## clarity^7     191.11     41.40   4.616 3.92e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3870 on 53926 degrees of freedom
## Multiple R-squared:  0.05937, Adjusted R-squared:  0.05915
## F-statistic: 261.8 on 13 and 53926 DF, p-value: < 2.2e-16
```

```
summary(without_inter_diamond_price_model)
```

```
##
## Call:
## lm(formula = price ~ 0 + color + clarity, data = diamonds)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -6046  -2469  -1308   1164  16858
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## colorD      2747.66      52.12  52.722 < 2e-16 ***
## colorE      2757.08      44.20  62.375 < 2e-16 ***
## colorF      3462.31      43.77  79.097 < 2e-16 ***
## colorG      3841.91      40.49  94.887 < 2e-16 ***
## colorH      4167.61      46.44  89.734 < 2e-16 ***
## colorI      4780.83      56.10  85.225 < 2e-16 ***
## colorJ      4933.65      76.03  64.890 < 2e-16 ***
## clarity.L   -1718.62     97.41 -17.642 < 2e-16 ***
## clarity.Q    -594.15     95.30  -6.235 4.56e-10 ***
## clarity.C     681.36     81.96   8.313 < 2e-16 ***
## clarity^4    -248.60     65.71  -3.783 0.000155 ***
## clarity^5     806.88     53.76  15.010 < 2e-16 ***
## clarity^6    -228.71     46.90  -4.876 1.08e-06 ***
## clarity^7     191.11     41.40   4.616 3.92e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3870 on 53926 degrees of freedom
## Multiple R-squared:  0.523, Adjusted R-squared:  0.5228
## F-statistic: 4223 on 14 and 53926 DF, p-value: < 2.2e-16
```

Which coefficients did not change between the models and why?

The clarity since they are numeric values.

Create a 2x2 matrix with the first column 1's and the next column iid normals. Find the absolute value of the angle (in degrees, not radians) between the two columns.

```
matrix_two = matrix(NA,2,2)
matrix_two[,1] = 1
matrix_two[,2] = rnorm(1)
theta_in_rad = acos(matrix_two[,1] %*% matrix_two[,2] / sqrt(sum(matrix_two[,1]^2) * sum(matrix_two[,2]^2)))
theta_in_rad * 180 / pi #Getting 0 or 180

##           [,1]
## [1,] 180
```

Repeat this exercise  $Nsim = 1e5$  times and report the average absolute angle.

```
theta_sum = 0
for(i in 1:1e5){
  theta_sum = theta_sum + acos(matrix_two[,1] %*% matrix_two[,2] / sqrt(sum(matrix_two[,1]^2) * sum(matrix_two[,2]^2)))
}
(theta_sum/1e5)* 180 / pi

##           [,1]
## [1,] 180
```

Create a 2xn matrix with the first column 1's and the next column iid normals. Find the absolute value of the angle (in degrees, not radians) between the two columns. For  $n \in 10, 50, 100, 200, 500, 1000$ , report the average absolute angle over  $Nsim = 1e5$  simulations.

```
matrix_three = matrix(NA,2,10)
matrix_three[,1] = 1
matrix_three[,1:10] = rnorm(10)
theta_in_rad = acos(matrix_three[,1] %*% matrix_three[,2] / sqrt(sum(matrix_three[,1]^2) * sum(matrix_three[,2]^2)))
theta_in_rad * 180 / pi

##           [,1]
## [1,] 96.79672

for(i in 1:1e5){
  theta_sum = theta_sum + acos(matrix_three[,1] %*% matrix_three[,2] / sqrt(sum(matrix_three[,1]^2) * sum(matrix_three[,2]^2)))
}
(theta_sum/1e5)* 180 / pi

##           [,1]
## [1,] 276.7967
```

What is this absolute angle converging to? Why does this make sense?

#TO-DO

Create a vector  $y$  by simulating  $n = 100$  standard iid normals. Create a matrix of size 100 x 2 and populate the first column by all ones (for the intercept) and the second column by 100 standard iid normals. Find the  $R^2$  of an OLS regression of  $y \sim X$ . Use matrix algebra.

```
y = rnorm(100)
X = matrix(NA,100,2)
X[,1] = 1
X[,2] = rnorm(100)
```

```

b = solve(t(X) %*% X ) %*% t(X) %*% y
y_hat = X %*% b
summary(y_hat)

##           V1
##  Min.      :-0.06150
##  1st Qu.: 0.02066
##  Median : 0.04091
##  Mean   : 0.04500
##  3rd Qu.: 0.07404
##  Max.    : 0.13723

SSE = sum((y - y_hat)^2)
SST = sum((y - mean(y))^2)
SSR = sum((y_hat - mean(y))^2)
1 - (SSE / SST)


## [1] 0.00161014

model = lm(y ~ X)
summary(model)


##
## Call:
## lm(formula = y ~ X)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2897 -0.7769  0.0414  0.6215  2.3976
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.05046    0.10935   0.461   0.646
## X1           NA         NA      NA     NA
## X2           0.04392    0.11049   0.398   0.692
##
## Residual standard error: 1.085 on 98 degrees of freedom
## Multiple R-squared:  0.00161,    Adjusted R-squared:  -0.008578
## F-statistic: 0.158 on 1 and 98 DF,  p-value: 0.6918

simple_df = data.frame(x = X[,2], y = y)
simple_viz_obj = ggplot(simple_df, aes(x, y)) +
  geom_point(size = 2)
simple_viz_obj

```



```
b_0 = model$coefficients[1]
b_1 = model$coefficients[3]
simple_ls_regression_line = geom_abline(intercept = b_0, slope = b_1, color = "red")
simple_viz_obj + simple_ls_regression_line
```



Write a for loop to each time bind a new column of 100 standard iid normals to the matrix  $X$  and find the  $R^2$  each time until the number of columns is 100. Create a vector to save all  $R^2$ 's. What happened??

```
X = matrix(1,100,1)
for(i in 2:100){
  X = cbind(X,rnorm(100))
}
b = solve(t(X) %*% X ) %*% t(X) %*% y
y_hat = X %*% b

SSE = sum((y - y_hat)^2)
SST = sum((y - mean(y))^2)
SSR = sum((y_hat - mean(y))^2)
1 - (SSE / SST) # It becomes just 100%
```

```
## [1] 1
```

```
dim(X)
```

```
## [1] 100 100
```

Add one final column to  $X$  to bring the number of columns to 101. Then try to compute  $R^2$ . What happens?

```
y = rnorm(101)
X = matrix(1,101,1)
for(i in 1:100){
  X = cbind(X,rnorm(100))
}
b = solve(t(X) %*% X ) %*% t(X) %*% y
```

```

## Error in solve.default(t(X) %*% X): system is computationally singular: reciprocal condition number :
y_hat = X %*% b

## Error in X %*% b: non-conformable arguments
SSE = sum((y - y_hat)^2)

## Error in eval(expr, envir, enclos): dims [product 100] do not match the length of object [101]
SST = sum((y - mean(y))^2)
SSR = sum((y_hat - mean(y))^2)
dim(X)

## [1] 101 101
1 - (SSE / SST) # It becomes just 100%

## [1] 1

```